

Normalized Floating Point Numbers

- We are most interested in normalized floating-point numbers, a format which includes:
 - sign
 - significand ($1.0 \leq \text{Significand} < \text{Radix}$)
 - integer power of the radix

Examples of Normalized Floating Point Numbers

These are normalized:

- $+1.23456789 \times 10^1$
- $-9.987654321 \times 10^{12}$
- $+5.0 \times 10^0$

These are **not** normalized:

- $+11.3 \times 10^3$ *significand > radix*
- -0.0002×10^7 *significand < 1.0*
- $-4.0 \times 10^{1/2}$ *exponent not integer*

Converting From Binary To Decimal

$$\begin{aligned}\mathbf{1.00101}_2 &= 1 \times 2^0 + 0 \times 2^{-1} + 0 \times 2^{-2} + 1 \times 2^{-3} \\ &\quad + 0 \times 2^{-4} + 1 \times 2^{-5} \\ &= 1 + 0/2 + 0/4 + 1/8 + 0/16 + 1/32 \\ &= 1 + 0.125 + 0.03125 \\ &= \mathbf{1.5625} \\ &= 37/32 = 1.5625\end{aligned}$$

Converting From Decimal To Binary

Let's start with $3.4625 \times 10^1 = 34.625$

Let's deal separately with the 34 (which equals 100010_2)

$$2 \times .625 = 1.25 \text{ (save the integer part)}$$

$$2 \times .25 = 0.5 \text{ (no integer part to save)}$$

$$2 \times .50 = 1.00 \text{ (save the integer part)}$$

Let's write them left to right in order:

$$34.625_{10} = 100010.101_2$$

Converting From Decimal To Binary – Another Example

$$1.23125 \times 10^1 = 12.3125$$

$$12_{10} = 1100_2$$

$$2 \times .3125 = \mathbf{0.625}$$

$$2 \times .625 = \mathbf{1.25}$$

$$2 \times .25 = \mathbf{0.50}$$

$$2 \times .50 = \mathbf{1.0}$$

$$\mathbf{12.3125}_{10} = \mathbf{1100.0101}_2$$

Normalizing Floating Point Data

Floating point data is normalized so that there is the significand is always one:

$$100001.101_2 = 1.00001101 \times 2^5$$

$$1100.0101_2 = 1.1000101 \times 2^3$$

Since the most significant bit is always 1, we can assume that it is implied and that we do not actually have to represent it.

Biased Exponents


- Short floating point numbers uses 8-bits for the exponents, which we want to range from -128 to +127.
- A biased exponent uses some value other than 0 as the baseline, which must be subtracted to get the actual exponent value.
- Example (in short floating point):
 - exponent 135 = $135 - 127 = 2^8$
 - exponent 120 = $120 - 127 = 2^{-7}$

Representing Floating Point Values In Memory

There are three standard formats for representing floating-point numbers:

- 32-bit format (*single-precision*)
- 64-bit format (*double-precision*)
- 80-bit format (*extended precision*)

Short Floating Point Numbers



Representing Values

$$-12.4375_{10} = -1100.0111_2$$

Short:

$$-1.10001110000...0000_2 \times 2^{3+127}$$


The diagram shows the conversion of the binary value $-1.10001110000...0000_2$ into the IEEE 754 Short format. The sign bit is 1 (negative). The biased exponent is 127 (01111111). The fraction is 10001110000...0000.

1 10000010 10001110000...0000

1100 0001 0100 0111 0000 ... 0000₂

= C1470000h

Long Floating Point Numbers



Representing Values

$$-12.4375_{10} = -1100.0111_2$$

Long:


$$-1.10001110000...0000_2 \times 2^3 + 1023$$

1 10000000010 10001110000...0000

1100 0000 0010 1000 1110 0000 ... 0000₂

= C028E00000000000h

Extended Floating Point Numbers



Representing Values

$$-12.4375_{10} = -1100.0111_2$$

Extended: $-1.10001110000...0000_2 \times 2^3$ ⁺¹⁶³⁸³

1 1000000000000010 10001110000...0000

1100 0000 0000 0010 1100 0111 0000 ... 0000₂
 = C002C70000000000000000h

Specifying Floating Point Data In Assembly Language

- We can use the:
 - **dd** (define doubleword) directive to allocate storage for single-precision floats
 - **dq** (define quadword) to allocate storage for double-precision floats and
 - **dt** (define tenbyte) for extended-precision floats.

Specifying Floating Point Data - An Example

- Allocating storage and initializing values

```
ShortOne  dd    1.0
LongOne   dq    1.0
Pi        dd    0.314159265E1
IntRate   dt    13.25E-1
```

Allocating storage without initializing:

```
Mass      dd    ?
CoefFric  dq    ?
Temp      dt    ?
```


Floating Point Operations

- Floating point operations include:
 - moving and rounding data
 - conversion
 - addition
 - subtraction
 - multiplication
 - division
 - remainder
 - comparison

Moving Floating Point Data

- Moving floating point data can be done using the standard **mov** instruction in Assembly language.
- If the source and destination are different length, care must be taken in conversion to ensure that exponent and significand are properly converted.

Data Conversion

- Integer and floating point data cannot be used interchangeably; data conversion is necessary and real-to-integer conversion is not without potential problems:
 - Underflow – a magnitude too small to represent as an integer.
 - Overflow – a magnitude too small to represent as an integer.
 - Inexact result – a loss of all of part of the fractional part of the floating-point fvalue.

Floating Point Addition

- To add two floating point values, they have to be aligned so that they have the same exponent.
- After addition, the sum may need to be normalized.
- Potential errors include overflow, underflow and inexact results.
- Examples:

2.34×10^3	6.22×10^8
$+ 0.88 \times 10^3$	$+ 3.93 \times 10^8$
3.22×10^3	$10.15 \times 10^8 = 1.015 \times 10^9$

Floating Point Subtraction

- Subtracting floating point values also requires re-alignment so that they have the same exponent.
- After subtraction, the difference may need to be normalized.
- Potential errors include overflow, underflow and inexact results, and the difference may have one significant bit less than the operands..
- Examples:

$$2.34 \times 10^3$$

$$\underline{-0.88 \times 10^3}$$

$$1.46 \times 10^3$$

$$6.44 \times 10^4$$

$$\underline{- 6.23 \times 10^4}$$

$$0.21 \times 10^4 = 2.1 \times 10^3$$

Floating Point Multiplication

- Multiplying floating point values does not requires re-alignment - realigning may lead to loss of significance.
- After multiplication, the product may need to be normalized.
- Potential errors include overflow, underflow and inexact results.
- Examples:

$$2.4 \times 10^{-3}$$

$$\times \underline{6.3 \times 10^2}$$

$$15.12 \times 10^1 = 1.512 \times 10^2$$

Floating Point Division

- Dividing floating point values does not require re-alignment.
- After division, the (floating point) quotient may need to be normalized – there is no remainder
- Potential errors include overflow, underflow, inexact results and attempts to divide by zero.
- Examples:

$$1.86 \times 10^{13} \div 7.44 \times 10^5 = \begin{array}{l} 0.25 \times 10^8 \\ 2.5 \times 10^7 \end{array}$$

Floating Point Remainder

- There is usually no remainder in floating point division, because the quotient can be a floating point value itself.
- Sometimes we want a remainder, i.e., the difference between the dividend and the product of the quotient rounded to the nearest integer) and the divisor:
- $s \text{ REM } t = s - t \times \text{NINT}(s/t)$
- Remainder will not produce inexact results, underflow or overflow but can lead to an attempt to divide by zero.

Floating Point Comparison

- There are usually three results that can happen as a result of floating point comparison:
 - less than
 - greater than
 - equal to
- In some instances, there is a fourth result: unordered, which occurs if one of the values is the result of an arithmetic error.
- These errors can result from adding or subtracting infinite values and are called *NaNs* (for *Not a Number*).

The Intel Floating Point Co-processors

- Early Intel processors (8088/8086, 80286, 80386) had no floating point capabilities; unless you wished to emulate floating point operations using software routines, you needed to add a co-processor (8087, 80287, 80387).
- 80486 and Pentium family processors include a floating point unit with an architecture that is the same as the coprocessors.