

# CONCURRENT AND PARALLEL PROGRAMMING

---

Dr. Lavika Goel  
MNIT Jaipur

# Syllabus - CST 303

- Concurrent versus sequential programming.
- Concurrent programming constructs and race condition.
- Synchronisation primitives.
- Processes and threads.
- Interprocess communication.
- Livelock and deadlocks, starvation, and deadlock prevention.
- Issues and challenges in concurrent programming paradigm and current trends.

# Syllabus - CST 303

- Parallel algorithms – sorting, ranking, searching, traversals, prefix sum etc.
- Parallel programming paradigms – Data parallel, Task parallel, Shared memory and message passing,
- Parallel Architectures, GPGPU, pthreads, STM, OpenMP
- OpenCL, Cilk++, Intel TBB, CUDA
- Heterogeneous Computing: C++AMP, OpenCL

# References

- Principles of Concurrent and Distributed Programming by Ben-Ari (Prentice-Hall International)
- Concurrent Programming: Principles and Practice by Greg Andrews (Addison Wesley)
- Synchronization Algorithms and Concurrent Programming by Gadi Taubenfeld (Pearson)

# References

- Introduction to Parallel Computing by Ananth Grama, et al (Pearson)
- Programming Massively Parallel Processors - A Hands-on Approach by David B. Kirk (Morgan Kaufmann)
- Parallel Algorithms by Joseph JaJa (Addison Wesley)
- CUDA Programming by Shane Cook (Morgan Kaufmann)
- Heterogeneous Computing with OpenCL by Benedict Gaster, et al (Morgan Kaufmann)

# Approach

- Concurrent Programming (POSIX Threads)
- Synchronization Primitives
- Interprocess Communication
- Livelock and Deadlocks
- Parallel Programming Paradigms
  - Message Passing (MPI), Shared Memory (OpenMP)
- Parallel Architectures GPGPU
- Nvidia CUDA, AMD OpenCL, Cilk++, Intel TBB
- Parallel Algorithms
- Heterogeneous Computing

# Weightage – Theory

- Mid Term 30
- End Term 50
- Quiz / Assignment 20

# Concurrent versus Parallel

- **Concurrency** is when two tasks can start, run, and complete in overlapping time periods.
- **Parallelism** is when tasks literally run at the same time, eg. on a multicore processor.
- **Concurrency:** A condition that exists when at least two threads are making progress.
- **Parallelism:** A condition that arises when at least two threads are executing simultaneously.




# Concurrent versus Parallel

- **Concurrency:** When multiple tasks performed simultaneously with shared resources
- **Parallelism:** when single task divided into multiple simple independent tasks which can be performed simultaneously
- **Concurrency:** Programming as the composition of independently executing processes.
- **Parallelism:** Programming as the simultaneous execution of (possibly related) computations.


# Concurrent versus Parallel

- Sequential
- Concurrent
- Parallel
- Concurrent But Not Parallel
- Parallel But Not Concurrent
- Concurrent and Parallel



# Problem !






# Solution ?




# Parallel ?






# Concurrent ?



# More Concurrent ?







# Parallel ?







# Concurrent or Parallel ?




# Concurrent or Parallel ?



# Multi Gopher + Staging Pile



# Full Optimization



# CONCURRENT PROGRAMMING

## CONCEPTS & NOTATIONS

---

# Three Issues

- How to express concurrent execution
- How processes communicate
- How processes synchronize

# Concurrent Programs

- Sequential program specifies sequential execution of a list of statements
- Execution of the program is called a process
- Concurrent program specifies two or more sequential programs that may be executed concurrently as parallel processes

# Concurrent Programs

- Concurrent program can be:
  - Executed either by allowing processes to share one or more processes – multiprogramming
  - Executed by running each process on its own processor
    - multiprocessing if the processors share common memory as in multiprocessor
    - distributed processing if processors are connected by a communication network



# Process Interaction

- Concurrently executing processors must communicate and synchronize in order to cooperate
- Communication allows execution of one process to influence execution of another
- Interprocess communication can be done
  - Use of shared variables
  - By message passing

# Synchronization

- Necessary when processes communicate
- Synchronization is a set of constraints on the ordering of events
- Programmer employs a synchronization mechanism to delay execution of a process in order to satisfy such constraints

# Concurrent Execution

- fork
- join
- cobegin
- call
- resume
- process

# Sync Prim – Shared Variables

- Busy Waiting
- Semaphores
  - Binary
  - Counting
- Conditional Critical Regions- mutual exclusion condition
- Monitors

# Sync Prim – Message Passing


- Specifying channels of Communication
  - send
  - receive
  - blocking
  - non blocking
  - asynchronous message
  - buffered
  - send no-wait

# Concurrent versus Parallel

- Concurrent program is a set of sequential programs that can be executed in parallel
- Parallel program is used for systems in which executions of several programs overlap in time by running them on separate processors
- Concurrency can be used to denote potential parallelism – execution may, but need not overlap


# Concurrency as Abstract Parallelism

- Clock speed of CPU of a PC is of the order of magnitude of one GHz (one billion times a second)
- Every nano second, the clock ticks and the processor performs some operation



# Concurrency as Abstract Parallelism

- If we are processing characters by hand, we do not consciously perform operations on the scale of nanoseconds
- Tremendous gap between human and electronic devices led to OS which allow I/O operations in parallel with computation

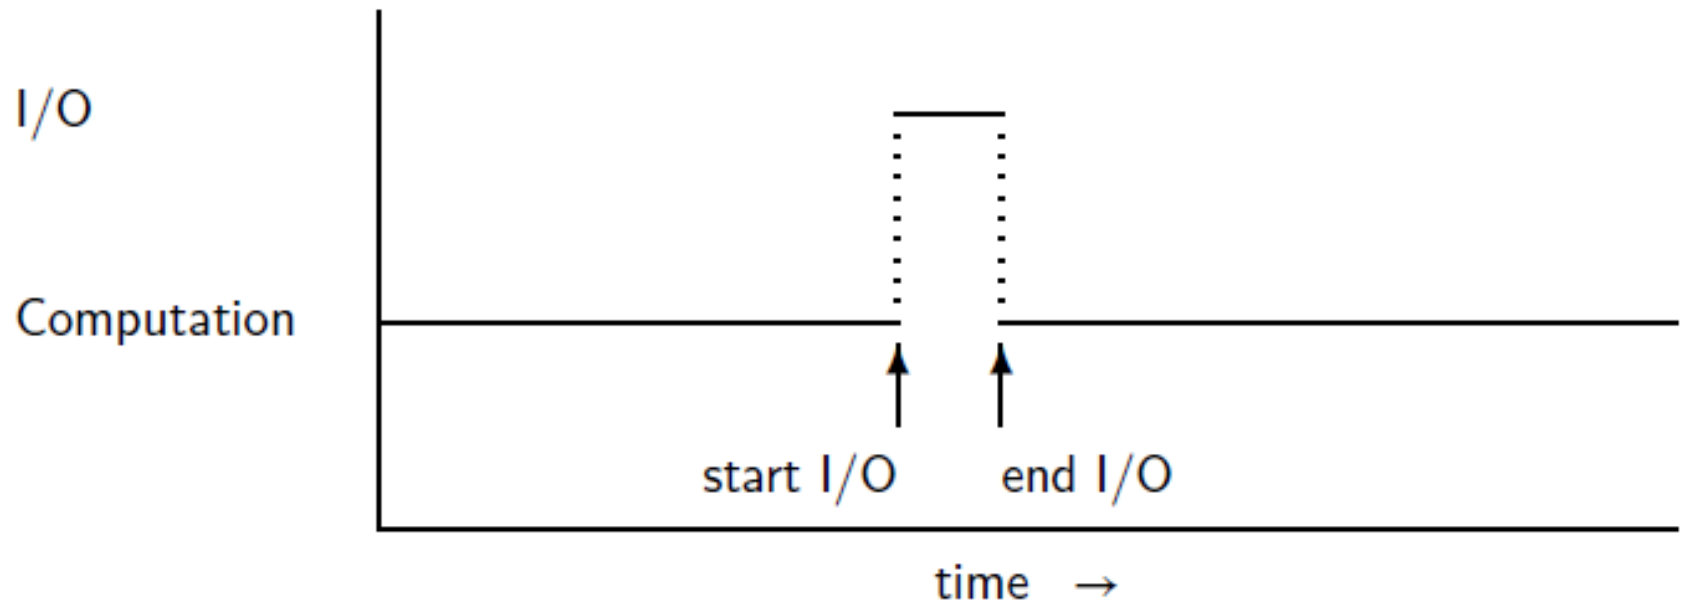




# Concurrency as Abstract Parallelism

- On a single CPU, processing required for each character typed cannot really be done in parallel
- A fraction of micro second can be stolen from the other computation to process the character
- Degradation is not noticeable even if the overhead of switching between the two computations is included
- I/O processing is performed as the result of an interrupt and is a separate process executed concurrently with a process doing another computation

# Concurrency as Abstract Parallelism



# Multitasking

- Simple generalization of overlapping the computation with that of another computation
- Kernel of all modern OS does *multitasking*
- Scheduler program is run by the OS to determine which process should be allowed to run for the next interval of time
- Time-slicing is implemented where computations are periodically interrupted to allow a fair sharing of resources of CPU

# Multitasking


- Programming languages support *multi threading* - multitasking within programs
- Threads enable the programmer to write concurrent computations within a single program
- Process runs in its own address space managed by the OS
- Thread runs within the address space of a single process which may be managed by a multithreaded kernel within the process

# Interleaving of Atomic Statements

- *Concurrent program* consists of a finite set of sequential processes
- The processes are written using a finite set of *atomic statements*
- Execution of a concurrent program proceeds by executing a sequence of atomic statements obtained by *arbitrarily interleaving* the atomic statements from the processes
- *Computation* is an execution sequence that can occur as a result of the interleaving

# Interleaving of Atomic Statements

- Computations are also called *scenarios*
- During a computation, the *control pointer* of a process indicates the next statement that can be executed by that process



# Concurrent and Serial Program

## Algorithm 2.1: Trivial concurrent program

integer  $n \leftarrow 0$

**p**

integer  $k1 \leftarrow 1$

p1:  $n \leftarrow k1$

**q**

integer  $k2 \leftarrow 2$

q1:  $n \leftarrow k2$

## Algorithm 2.2: Trivial sequential program

integer  $n \leftarrow 0$

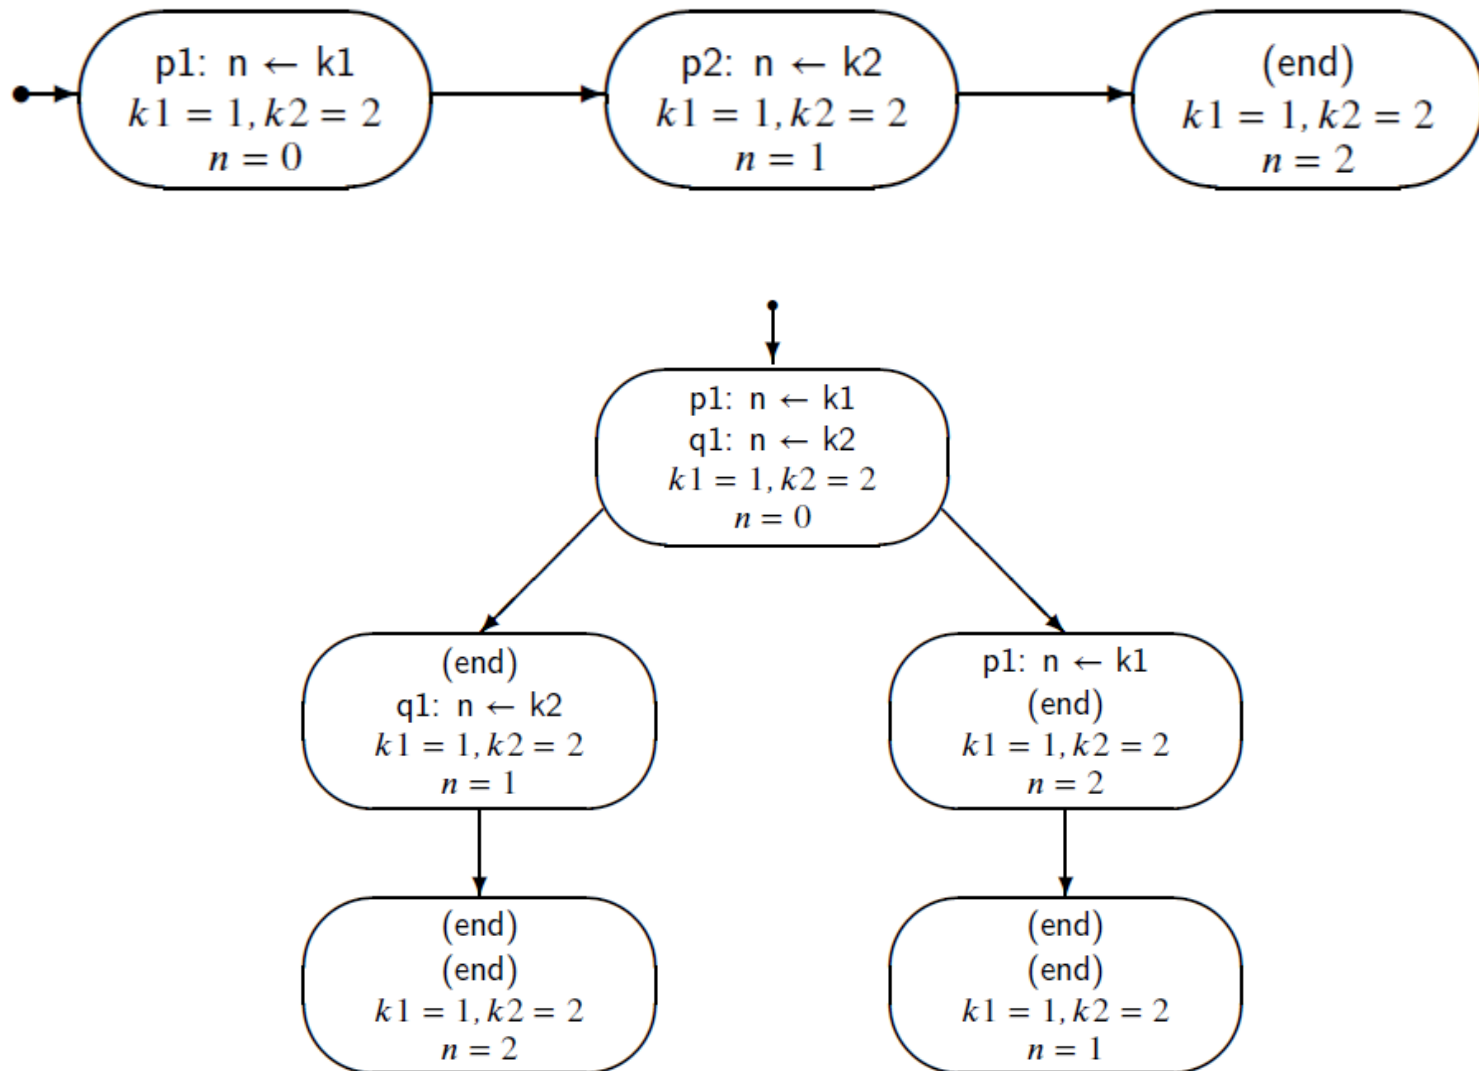
integer  $k1 \leftarrow 1$

integer  $k2 \leftarrow 2$

p1:  $n \leftarrow k1$

p2:  $n \leftarrow k2$

# State Representation





# State and Transition

- The *state* of a concurrent algorithm is a tuple consisting of one element for each process that is a label from that process, and one element for each global or local variable that is a value whose type is the same as type of the variable
- Let  $s_1$  and  $s_2$  be two states, then there is a *transition* between  $s_1$  and  $s_2$  if executing a statement in state  $s_1$  changes the state to  $s_2$

# State Diagram

- A *state diagram* is a graph defined inductively.
- The initial state diagram consists a single node labeled with the initial state.
- If state  $s_1$  labels a node in the state diagram, and if there is a transition from  $s_1$  to  $s_2$ , then there is a node labeled  $s_2$  in the state diagram and a directed edge from  $s_1$  to  $s_2$
- For each state, there is only one node labeled with that state
- Set of reachable states is the set of states in a state diagram

# Scenario

- Scenario is defined by a sequence of states
- Use tables of scenarios instead of diagrams
- List the sequence of states, columns for control pointers are labeled with processes and columns for variable values with the variable names

Process p	Process q	n	k1	k2
<b>p1: n ← k1</b>	<b>q1: n ← k2</b>	0	1	2
(end)	<b>q1: n ← k2</b>	1	1	2
(end)	(end)	2	1	2

# Scenario

- In a state, there may be more than one statement that can be executed
- Bold font is used for denoting the statement that was executed to get to the state in the following row
- Rows represent states
- If a statement executed is an assignment statement, the new value that is assigned to the variable is a component of the next state in the scenario, which is found in the next row.


# Justification of the Abstraction

- Coordination between computers cannot be done at the level of individual instructions
- There is not enough time to coordinate individual instructions of more than one CPU
- Execution of a concurrent program is considered to be carried out by a global entity who at each step selects the process from which the next statement will be executed
- The interleaving of processes is arbitrary
- This abstraction is justified for various possible architectures


# State Diagram for the Frog Puzzle



# State Diagram for the Frog Puzzle

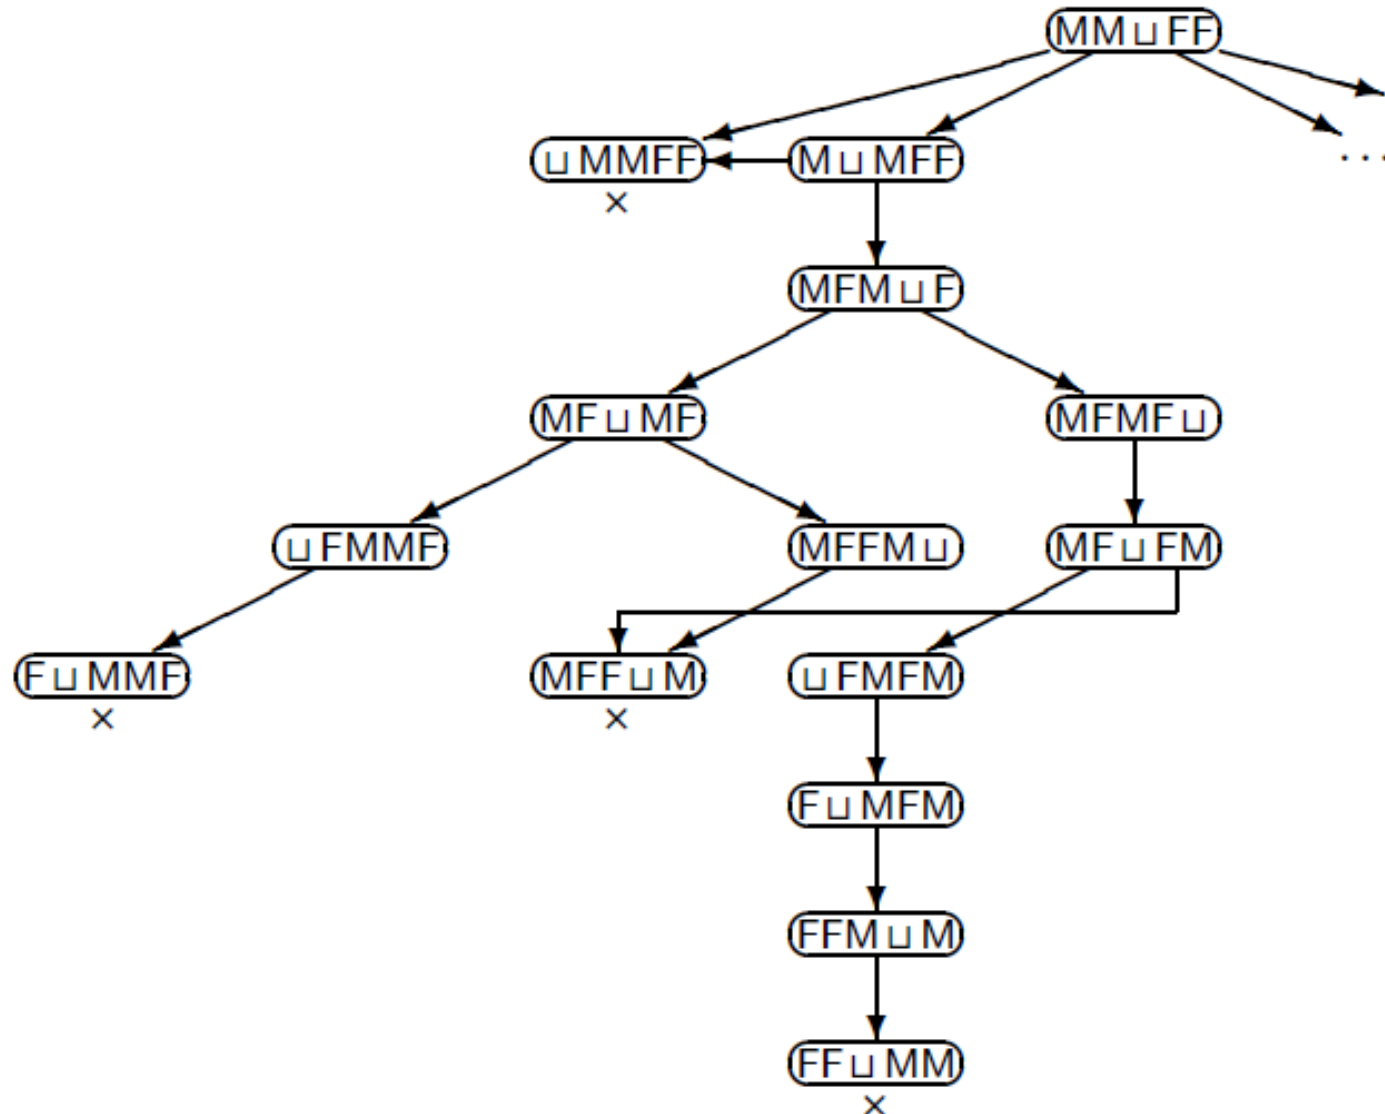


# State Diagram for the Frog Puzzle





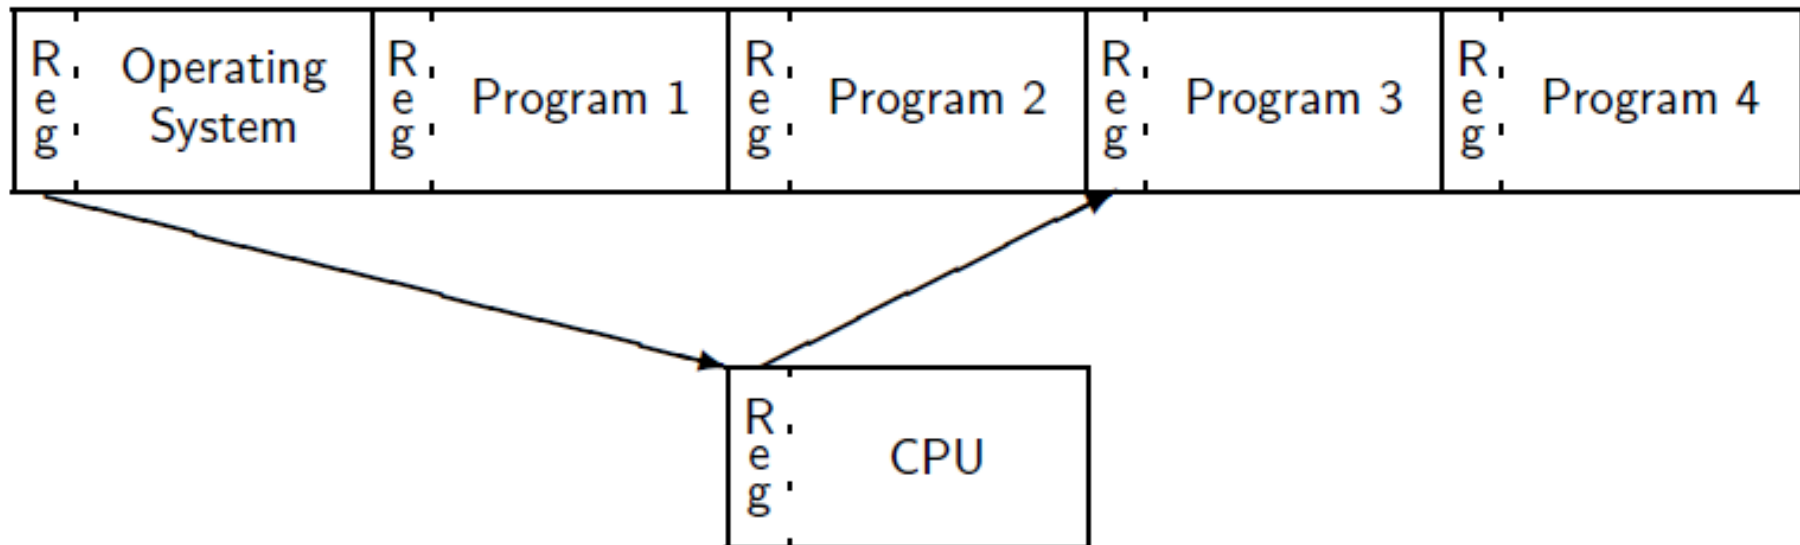
# State Diagram for the Frog Puzzle



# Multitasking Systems

- Selection of next instruction to execute is carried out by the CPU and the OS
- Interrupts from I/O devices or interrupt timers will cause execution to be interrupted
- *Interrupt handler* will be executed and upon its completion, the OS scheduler is invoked to select a new process
- This mechanism is called *context switch*

# Multitasking Systems

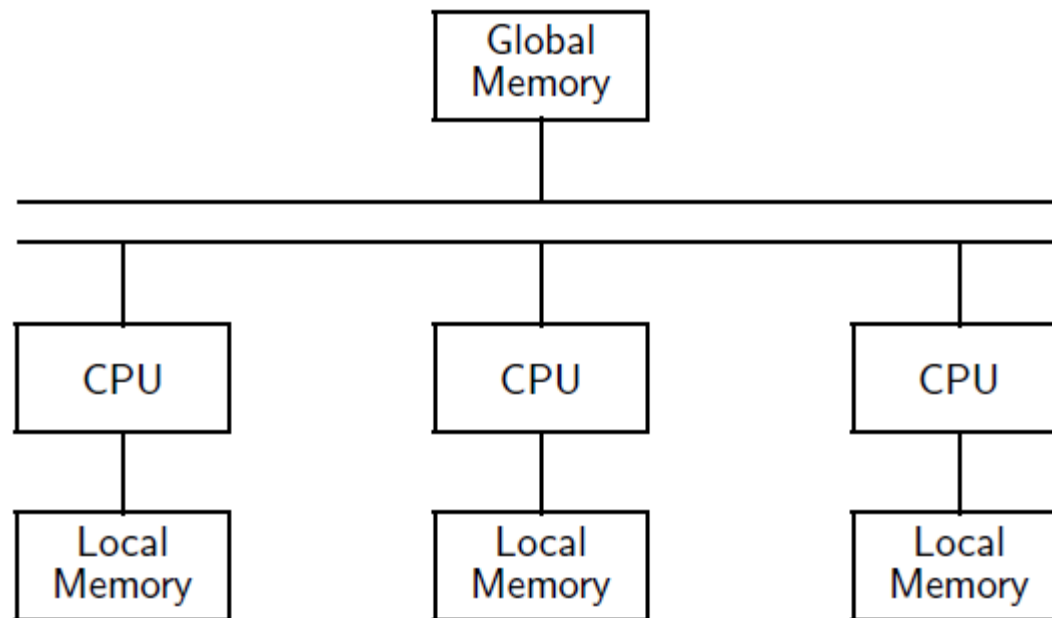


# Multitasking Systems

- When execution is interrupted, all the registers (computation, control and others) are saved in a predefined area in the program's memory
- Register contents required to execute the interrupt handler are loaded into the CPU
- After interrupt is processed, the symmetric context switch is performed storing the interrupt handler registers and loading program registers
- OS scheduler decides to perform the context switch with another program and not the one that was interrupted

# Multiprocessor Computers

- Computer with more than one CPU
- Memory is physically divided into banks
- *Local* memory (accessed by one CPU) and *Global* memory (accessed by all CPUs)




# Multiprocessor Computers

- If sufficient CPUs are present, each process has its own CPU
- Interleaving is no longer real as each CPU is executing instructions independently
- As long as there is no *contention* and two CPUs do not attempt to access the same resource (say global memory)
- Computations defined by interleaving will be indistinguishable from parallel execution
- With contention there is a potential problem

# Multiprocessor Computers

- If two processes try to read or write to a cell simultaneously so that the operations overlap



# Multiprocessor Computers



- Interleaving is handled as memory hardware is designed so that one access completes before the other commences
- If two CPUs attempt to read or write the same cell in memory, the result will be same as if the two instructions were executed in either order
- Atomicity and interleaving are performed by the hardware
- Arbitrary interleaving makes more sense as there is no central scheduler, any computation resulting from interleaving may occur



# Distributed Systems

- Distributed system is composed of several computers that have no global resources
- They are connected by communication *channels* enabling them to send messages to each other
- Abstraction of interleaving is not applicable
- Each node is either executing one of its statements, sending a message or receiving a message
- Any interleaving of all the events of all the nodes can be used for reasoning
- Interleaving is consistent with the statement sequences of individual node and with the requirement that message be sent before it is received

# Distributed Systems



# Distributed vs Concurrent

- In concurrent systems, global memory is accessible to all systems
- In distributed systems, nodes may be geographically distant from one another.
- Each node may not send a message *directly* to another node
- We have to consider the *topology* or *connectedness* of the system – fully or ring
- Hardware / software failures catastrophic in CS
- In DS, failures are catastrophic for single nodes, but faulty nodes can be diagnosed and messages can be relayed thro alternate paths

# Atomic Statements

- Concurrent programming abstraction has been defined in terms of interleaving of atomic statements
- An atomic statement is executed to completion without the possibility of interleaving statements from another process
- An important property of atomic statements is that if two are executed simultaneously the result is the same as if they had been executed sequentially (in either order)
- It is important to specify the atomic statements precisely as the correctness of the algorithm depends on the specification

# Atomic Statements

## Algorithm 2.3: Atomic assignment statements

integer  $n \leftarrow 0$

**p**

**q**

**p1:**  $n \leftarrow n + 1$

**q1:**  $n \leftarrow n + 1$

Process p	Process q	n
<b>p1:</b> $n \leftarrow n + 1$	<b>q1:</b> $n \leftarrow n + 1$	0
(end)	<b>q1:</b> $n \leftarrow n + 1$	1
(end)	(end)	2

Process p	Process q	n
<b>p1:</b> $n \leftarrow n + 1$	<b>q1:</b> $n \leftarrow n + 1$	0
<b>p1:</b> $n \leftarrow n + 1$	(end)	1
(end)	(end)	2

# Atomic Statements with Global Reference

**Algorithm 2.4: Assignment statements with one global reference**

integer  $n \leftarrow 0$

**p**

**q**

integer temp

p1:  $\text{temp} \leftarrow n$

p2:  $n \leftarrow \text{temp} + 1$

integer temp

q1:  $\text{temp} \leftarrow n$

q2:  $n \leftarrow \text{temp} + 1$

Process p	Process q	n	p.temp	q.temp
<b>p1: temp←n</b>	q1: temp←n	0	?	?
<b>p2: n←temp+1</b>	q1: temp←n	0	0	?
(end)	<b>q1: temp←n</b>	1	0	?
(end)	<b>q2: n←temp+1</b>	1	0	1
(end)	(end)	2	0	1

# Atomic Statements with Global Reference

**Algorithm 2.4: Assignment statements with one global reference**

integer  $n \leftarrow 0$

**p**

**q**

integer temp

p1: temp  $\leftarrow$  n

p2: n  $\leftarrow$  temp + 1

integer temp

q1: temp  $\leftarrow$  n

q2: n  $\leftarrow$  temp + 1

Process p	Process q	n	p.temp	q.temp
<b>p1: temp</b> $\leftarrow$ n	q1: temp $\leftarrow$ n	0	?	?
p2: n $\leftarrow$ temp+1	<b>q1: temp</b> $\leftarrow$ n	0	0	?
<b>p2: n</b> $\leftarrow$ temp+1	q2: n $\leftarrow$ temp+1	0	0	0
(end)	<b>q2: n</b> $\leftarrow$ temp+1	1	0	0
(end)	(end)	1	0	0

# Correctness

- In sequential program, running a program with the same input will always give the same result
- In a concurrent program, some scenarios may give correct output while others don't
- Correctness of sequential programs is defined in terms of computing a final result
- Correctness of concurrent programs is defined in terms of properties of computation
  - There are two properties – *Safety and Liveness*



# Correctness

- **Safety:** The property must always be true
  - For a safety property  $P$  to hold, it must be true that in every state of every computation,  $P$  is true
- **Liveness:** The property must eventually become true
  - For a liveness property  $P$  to hold, it must be true that in every computation, there is some state in which  $P$  is true

# Correctness

- The challenge of concurrent programs is to ensure that they do useful things, satisfying liveness properties – without violating safety properties
- Safety and liveness are duals of each other – negation of safety property is liveness and v. v
- Safety property will be true if and only if the Liveness property will be false
- Linear Temporal Logic can be used (BTL) to demonstrate correctness

# Fairness

- A scenario is (weakly) *fair* if at any state in the scenario, a statement that is continually enabled eventually appears in the scenario
- Does this program halt?

Algorithm 2.5: Stop the loop A	
integer $n \leftarrow 0$ boolean $\text{flag} \leftarrow \text{false}$	
p	q
p1: while $\text{flag} = \text{false}$ p2: $n \leftarrow 1 - n$	q1: $\text{flag} \leftarrow \text{true}$ q2:

# Example

## Algorithm 2.9: Concurrent counting algorithm

integer  $n \leftarrow 0$

**p**

**q**

integer temp

integer temp

p1: do 10 times

q1: do 10 times

p2:     temp  $\leftarrow$  n

q2:     temp  $\leftarrow$  n

p3:     n  $\leftarrow$  temp + 1

q3:     n  $\leftarrow$  temp + 1

# C Code

```
int n = 0;
```

```
void p() {  
    int temp, i;  
    for (i = 0; i < 10; i++) {  
        temp = n;  
        n = temp + 1;  
    }  
}
```

```
void q() {  
    int temp, i;  
    for (i = 0; i < 10; i++) {  
        temp = n;  
        n = temp + 1;  
    }  
}
```

```
void main() {  
    cobegin { p(); q(); }  
    cout << "The value of n is " << n << "\n";  
}
```

# Java Code

```
class Count extends Thread {  
    static volatile int n = 0;  
  
    public void run() {  
        int temp;  
        for (int i = 0; i < 10; i++) {  
            temp = n;  
            n = temp + 1;  
        }  
    }  
}
```

```
        public static void main(String[] args) {  
            Count p = new Count();  
            Count q = new Count();  
            p.start ();  
            q.start ();  
            try {  
                p.join ();  
                q.join ();  
            }  
            catch (InterruptedException e) { }  
            System.out.println ("The value of n is " + n);  
        }  
    }  
}
```