

CRITICAL SECTION PROBLEM

Dr. Lavika Goel

Critical Section Problem

- Each of N processes is executing in an infinite loop, a sequence of statements that can be divided into two subsequences:
 - *critical section*
 - *non critical section*
- Three correctness specifications are required of any solution:
 - *Mutual Exclusion*
 - *Freedom from deadlock*
 - *Freedom from starvation*

Critical Section Problem

- Mutual Exclusion
 - Statements from the critical sections of two or more processes must not be interleaved
- Freedom from deadlock
 - If **some** processes are trying to enter their critical sections, then **one** of them must eventually succeed
- Freedom from starvation
 - If **any** process tries to enter its critical section, then **that** process must eventually succeed

Critical Section Problem

- A *synchronization* mechanism must be provided to ensure that the correctness requirements are met.
- Synchronization mechanism consists of additional statements that are placed before and after the critical section
- The statements placed before critical section are called *preprotocol* and those after it are called *postprotocol*

Critical Section Problem

Algorithm 3.1: Critical section problem

global variables

p

q

local variables

loop forever

non-critical section

preprotocol

critical section

postprotocol

local variables

loop forever

non-critical section

preprotocol


critical section

postprotocol

Critical Section Problem

- The protocol may require local or global variables
- Critical section must progress
 - Once a process starts to execute the statements in a critical section, it must eventually finish execution of those statements
- The Non-Critical section need not progress
 - If the control pointer of a process is in its non-critical section, the process may terminate or enter an infinite loop

Critical Section Problem



Correctness

Algorithm 3.3: History in a sequential algorithm

integer $a \leftarrow 1$, $b \leftarrow 2$

p1: Millions of statements

p2: $a \leftarrow (a+b)*5$

p3: ...

Algorithm 3.4: History in a concurrent algorithm

integer $a \leftarrow 1$, $b \leftarrow 2$

p

p1: Millions of statements

p2: $a \leftarrow (a+b)*5$

p3: ...

q

q1: Millions of statements

q2: $b \leftarrow (a+b)*5$

q3: ...

Correctness


- Sequential
 - $s_i = (p2, 10, 20)$ and $s_{i+1} = (p3, 150, 20)$
- Concurrent
 - $s_i = (p2, q2, 10, 20)$
 - $s^p_{i+1} = (p3, q2, 150, 20)$ or $s^q_{i+1} = (p2, q3, 10, 150)$
- The set of reachable states are the only states that can appear in any computation
- To check correctness, it is only necessary to examine the set of reachable states and the transitions among them

First Attempt

- await $\text{turn} = 1$ waits until the condition $\text{turn} = 1$ becomes true
- This can be implemented by a busy-wait loop

Algorithm 3.2: First attempt	
integer $\text{turn} \leftarrow 1$	
p	q
loop forever	loop forever
p1: non-critical section	q1: non-critical section
p2: await $\text{turn} = 1$	q2: await $\text{turn} = 2$
p3: critical section	q3: critical section
p4: $\text{turn} \leftarrow 2$	q4: $\text{turn} \leftarrow 1$

State Diagram – First steps



Algorithm 3.2: First attempt

integer turn \leftarrow 1

p

q

loop forever

loop forever

p1: non-critical section

q1: non-critical section

p2: await turn = 1

q2: await turn = 2

p3: critical section

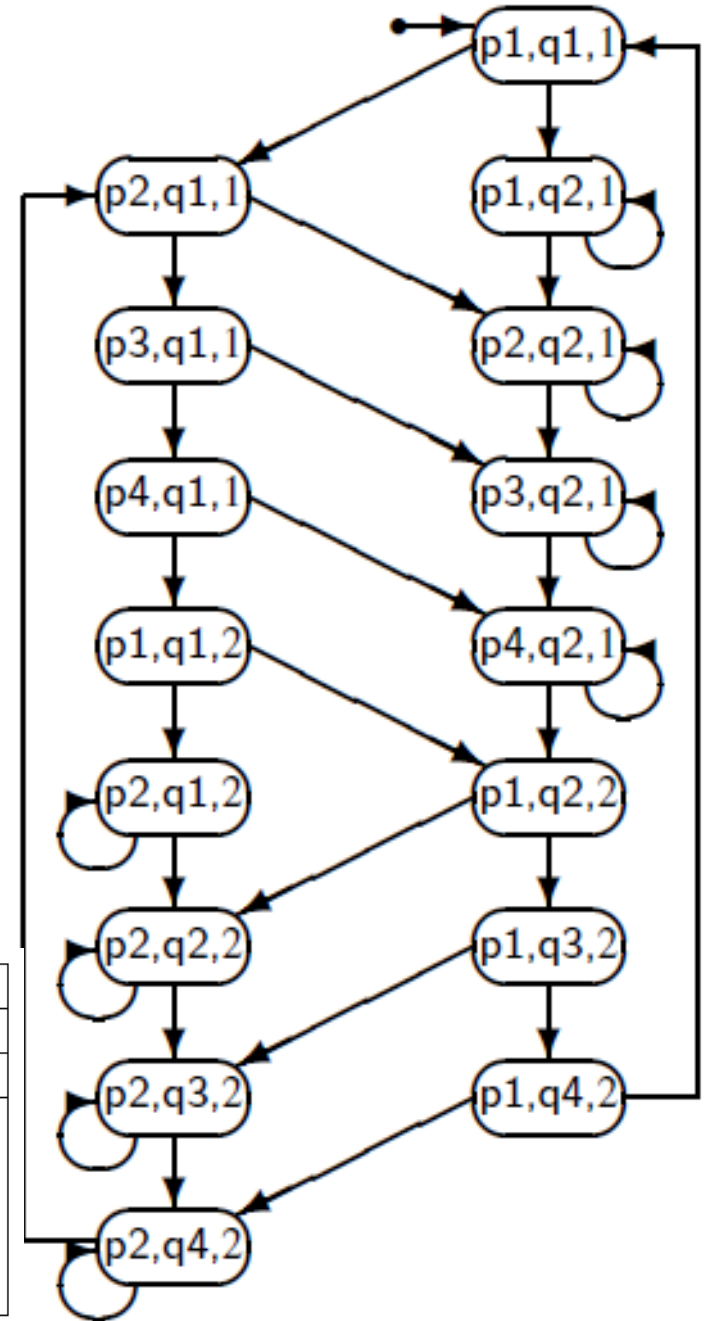
q3: critical section

p4: turn \leftarrow 2

q4: turn \leftarrow 1

Sixteen steps

- $(p_3, q_{3,1})$ or $(p_3, q_{3,2})$ do not occur – mutual exclusion property holds



Algorithm 3.2: First attempt	
integer turn $\leftarrow 1$	
p	q
loop forever	loop forever
p1: non-critical section	q1: non-critical section
p2: await turn = 1	q2: await turn = 2
p3: critical section	q3: critical section
p4: turn $\leftarrow 2$	q4: turn $\leftarrow 1$

First Attempt (Abbreviated)

Algorithm 3.5: First attempt (abbreviated)

integer turn \leftarrow 1

p

loop forever

p1: await turn = 1

p2: turn \leftarrow 2

q

loop forever

q1: await turn = 2

q2: turn \leftarrow 1

Algorithm 3.2: First attempt

integer turn \leftarrow 1

p

loop forever

p1: non-critical section

p2: await turn = 1

p3: critical section

p4: turn \leftarrow 2

q

loop forever

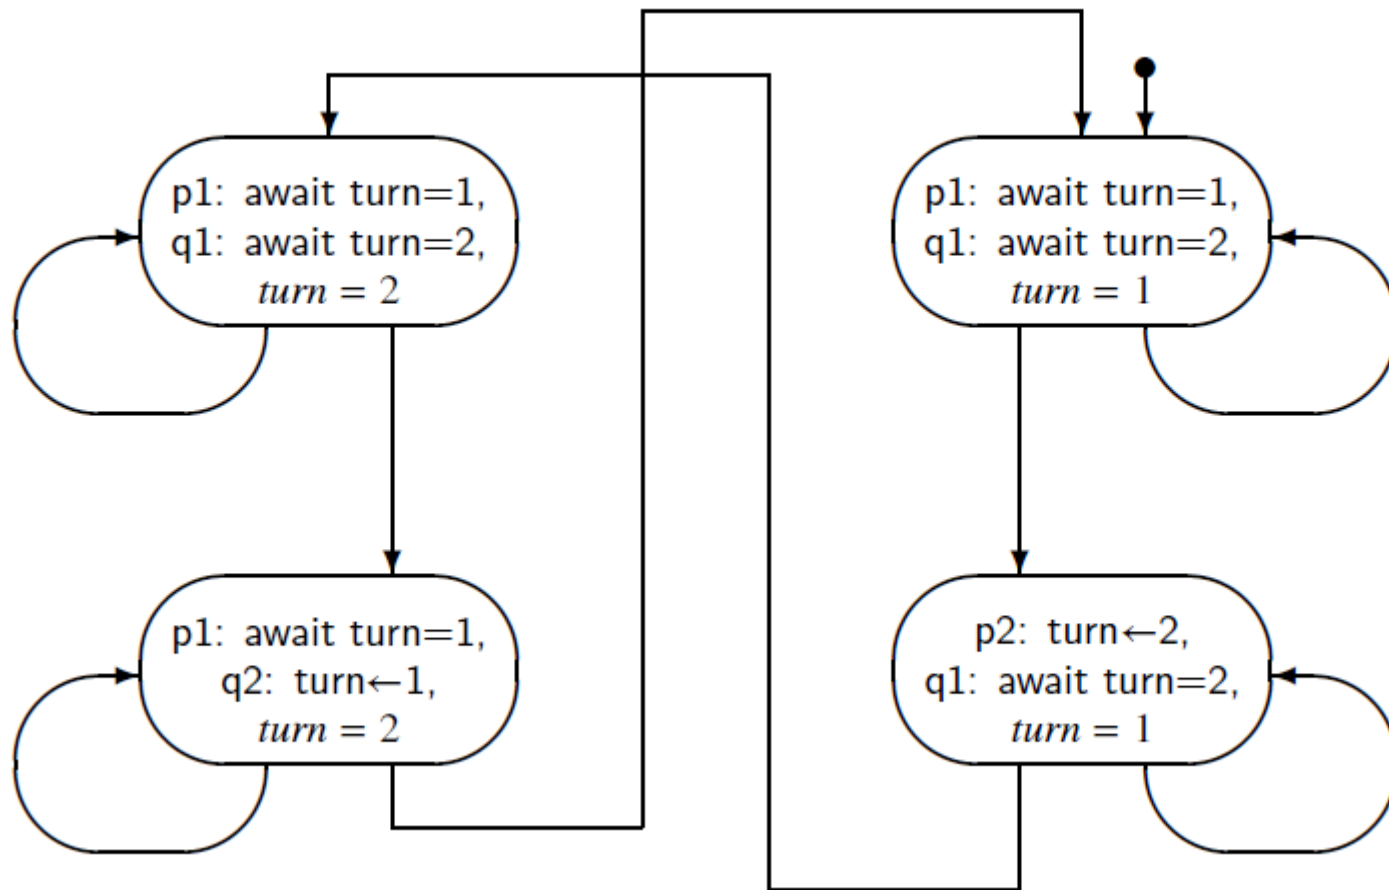
q1: non-critical section

q2: await turn = 2

q3: critical section

q4: turn \leftarrow 1

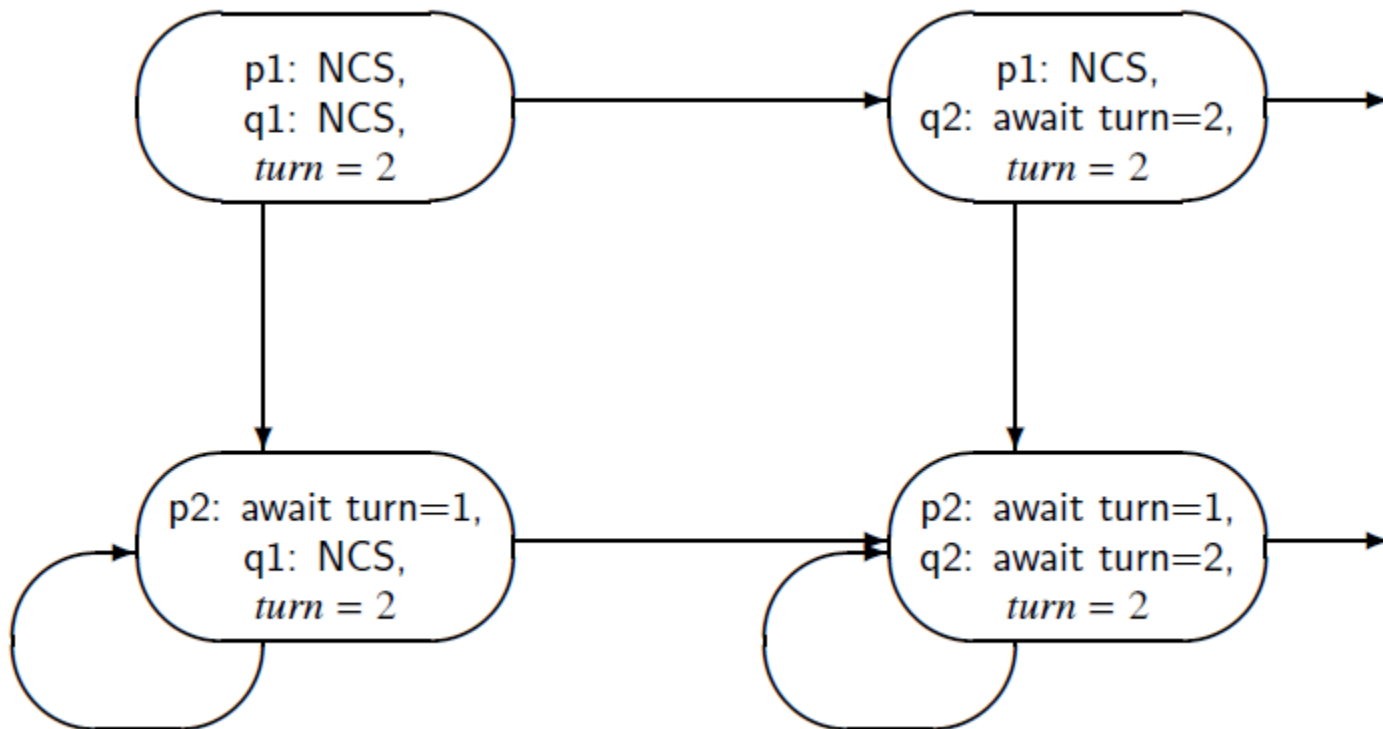
State Diagram – First Attempt (Abbv)



Correctness of First Attempt

- Proof of Mutual exclusion is immediate from the state diagram
- Proof of freedom from deadlock
 - If **some** processes are trying to enter their critical sections, then **one** of them must eventually succeed
- Proof of freedom from starvation
 - There is always some process holding the permission resource, so some process can always enter the CS ensuring there is no deadlock
 - If the process holding the permission resource remains indefinitely in its NCS, other process will never receive the resource and will never enter CS

State Diagram – Fragment



Second Attempt

- First attempt – both processes set and tested a single variable.
- If one process dies, other is blocked
- Each process is now given its own variable
- `want i` is true from step where process i wants to enter its critical section until it leaves
- `await` statements ensure that a process does not enter its CS while another process has its flag set

Second Attempt

Algorithm 3.6: Second attempt	
boolean wantp \leftarrow false, wantq \leftarrow false	
p	q
loop forever	loop forever
p1: non-critical section	q1: non-critical section
p2: await wantq = false	q2: await wantp = false
p3: wantp \leftarrow true	q3: wantq \leftarrow true
p4: critical section	q4: critical section
p5: wantp \leftarrow false	q5: wantq \leftarrow false

Second Attempt

- If a process halts in its critical section, the value of its variable `want` will remain false and the other process will always succeed in immediately entering the critical section
- Solves the problem of *starvation*
- But as we move ahead, we see that *mutual exclusion* property is not satisfied.

Second Attempt (Abbreviated)

Algorithm 3.7: Second attempt (abbreviated)

boolean wantp \leftarrow false, wantq \leftarrow false

p

loop forever

p1: await wantq = false

p2: wantp \leftarrow true

p3: wantp \leftarrow false

q

loop forever

q1: await wantp = false

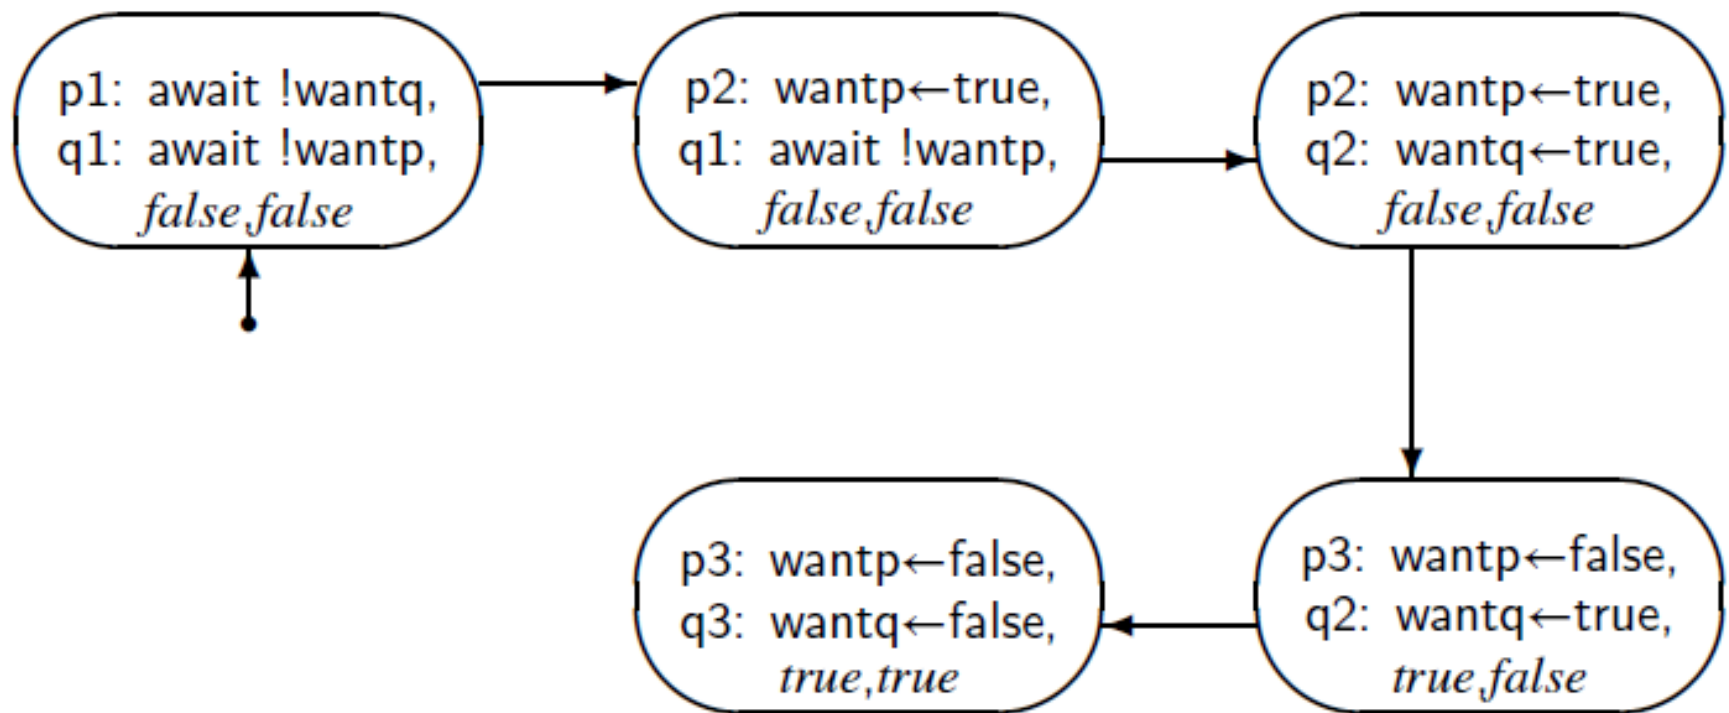
q2: wantq \leftarrow true

q3: wantq \leftarrow false

Tabular Form

Process p	Process q	wantp	wantq
p1: await wantq=false	q1: await wantp=false	<i>false</i>	<i>false</i>
p2: wantp←true	q1: await wantp=false	<i>false</i>	<i>false</i>
p2: wantp←true	q2: wantq←true	<i>false</i>	<i>false</i>
p3: wantp←false	q3: wantq←true	<i>true</i>	<i>false</i>
p3: wantp←false	q3: wantq←false	<i>true</i>	<i>true</i>

State Diagram – Fragment



Second Attempt

- To prove that mutual exclusion holds, it must be checked that no forbidden state appears in any scenario
- If mutual exclusion does in fact hold, we need to construct the full state diagram for the algorithm, because every path in the diagram is a scenario
- Every state must be examined to make sure it is not a forbidden state
- We can stop construction if a forbidden state is encountered.

Third Attempt

- Second attempt – variables `want` are intended to indicate when a process is in its critical section
- Once a process has completed its `await`, it cannot be prevented from entering its CS
- The state reached after `await` but before assignment to `want` is effectively part of CS, but value of `want` does not indicate this

Third Attempt

- Recognizes that the `await` statement should be part of the critical section by moving the assignment to `want` before the `await`
- The construction of the state diagram shows that mutual exclusion is not violated
- However, the algorithm can *deadlock* as shown

Third Attempt

Algorithm 3.8: Third attempt

boolean wantp \leftarrow false, wantq \leftarrow false

p	q
loop forever p1: non-critical section p2: wantp \leftarrow true p3: await wantq = false p4: critical section p5: wantp \leftarrow false	loop forever q1: non-critical section q2: wantq \leftarrow true q3: await wantp = false q4: critical section q5: wantq \leftarrow false

Algorithm 3.6: Second attempt

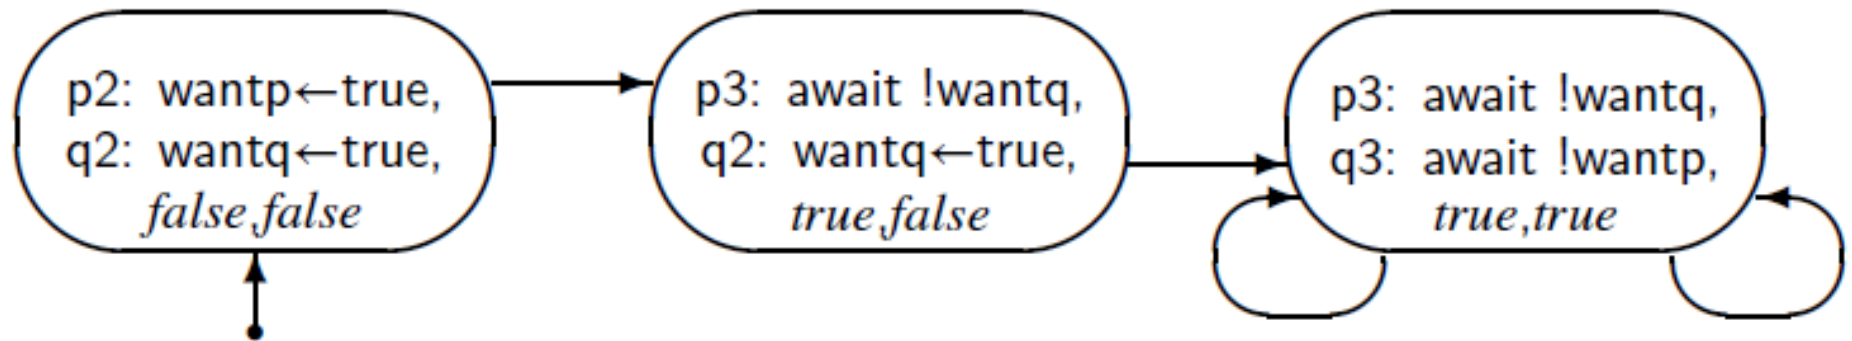
boolean wantp \leftarrow false, wantq \leftarrow false

p	q
loop forever p1: non-critical section p2: await wantq = false p3: wantp \leftarrow true p4: critical section p5: wantp \leftarrow false	loop forever q1: non-critical section q2: await wantp = false q3: wantq \leftarrow true q4: critical section q5: wantq \leftarrow false

Tabular Form

Process p	Process q	wantp	wantq
p1: non-critical section	q1: non-critical section	<i>false</i>	<i>false</i>
p2: wantp←true	q1: non-critical section	<i>false</i>	<i>false</i>
p2: wantp←true	q2: wantq←true	<i>false</i>	<i>false</i>
p3: await wantq=false	q2: wantq←true	<i>true</i>	<i>false</i>
p3: await wantq=false	q3: await wantp=false	<i>true</i>	<i>true</i>

State Diagram – Fragment



Deadlock and Livelock

- Term deadlock is usually used with a frozen computation where nothing whatsoever is being computed
- A scenario where several processes are actively executing statements, but nothing useful gets done is called livelock

Livelock

- Attempting to use resource preemption approaches to preventing deadlock may cause livelock: Processes don't deadlock, but fail to make progress either
- A thread often acts in response to the action of another thread. If the other thread's action is also a response to the action of another thread, then *livelock* may result.
- As with deadlock, livelocked threads are unable to make further progress. However, the threads are not blocked — they are simply too busy responding to each other to resume work.
- *Deadlock*: "Me first, Me first" *Livelock*: "You first, You first"

Livelock

- Wayne and Larry want to listen to a CD on CD player
- Larry has the CD but wants the CD player
- Wayne has the CD player but wants the CD
- Wayne steals (i.e. preempts) the CD from Larry, meanwhile Larry steals the CD player from Wayne
- Now, Wayne has the CD and Larry has the CD player
- Wayne steals the CD player from Larry, meanwhile Larry steals the CD from Wayne
- Now, Wayne has the CD player and Larry has the CD

...

Livelock

Deadlock vs Livelock

- **Deadlock**
 - It happens when a process waits for another one who is using some needed resource to finish with it, while the other process also wait for the first process to release some other resource.
- **Livelock**
 - A Livelock looks like a deadlock in the sense that two (or more) processes are blocking each others. But with the livelock, each process is waiting “actively”, trying to resolve the problem on its own (like reverting back its work and retry).

Deadlock vs Livelock

- A livelock is similar to a deadlock, except that the states of the processes involved in the livelock constantly change with regard to one another, none progressing.

Fourth Attempt

- Third attempt – when a process sets a variable want to be true, not only does it indicate its intention to enter its critical section, but also insists on its right to do so
- Deadlock occurs when both processes simultaneously insist on entering their CS
- Requires a process to give up its intention to enter the CS if it discovers that it is contending with the other process

Fourth Attempt

Algorithm 3.9: Fourth attempt

boolean wantp \leftarrow false, wantq \leftarrow false

p

loop forever

p1: non-critical section

p2: wantp \leftarrow true

p3: while wantq

p4: wantp \leftarrow false

p5: wantp \leftarrow true

p6: critical section

p7: wantp \leftarrow false

q

loop forever

q1: non-critical section

q2: wantq \leftarrow true

q3: while wantp

q4: wantq \leftarrow false

q5: wantq \leftarrow true

q6: critical section

q7: wantq \leftarrow false

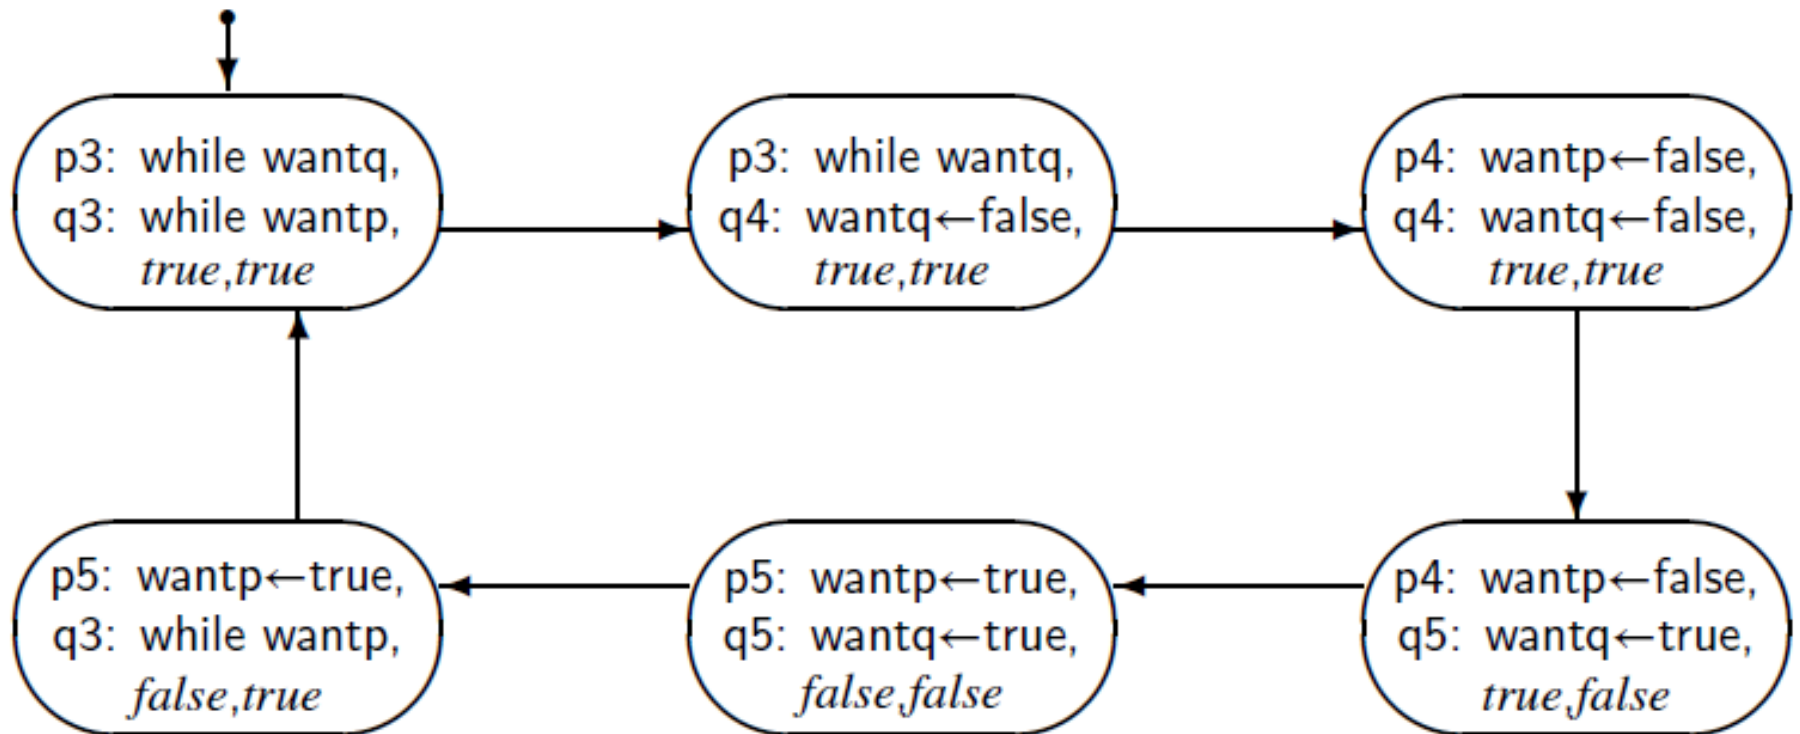
Fourth Attempt

- Statements p4 and p5 will be meaningless in a sequential algorithm, but useful in concurrent algorithm
- Arbitrary interleaving from and between the two processes, second process may execute an arbitrary number of statements between the two assignment statements to wantp
- When process p relinquishes the attempt to enter the critical section by resetting wantp to false, process q may now execute the wait statement and succeed in entering the CS

Fourth Attempt

- A state diagram will show that the mutual exclusion property holds and that there is no deadlock
- Scenario for starvation exists as shown in figure in next slide
- If interleaving is perfect, where the execution of a statement of process q is always followed by the execution of an equivalently numbered statement of process p, both the processes are starved

State Diagram – Fragment



Dekker's Algorithm

- Combination of first and fourth attempts
- First Attempt – explicitly passed the right to enter the Critical Section
- This caused processes to be closely coupled and prevented correct behavior in absence of contention
- Fourth Attempt – each process had its own variable which prevented problems in absence of contention
- Both processes however insist on entering CS in the presence of contention

Dekker's Algorithm

Algorithm 3.10: Dekker's algorithm

boolean wantp \leftarrow false, wantq \leftarrow false

integer turn \leftarrow 1

p

loop forever

p1: non-critical section

p2: wantp \leftarrow true

p3: while wantq

p4: if turn = 2

p5: wantp \leftarrow false

p6: await turn = 1

p7: wantp \leftarrow true

p8: critical section

p9: turn \leftarrow 2

p10: wantp \leftarrow false

q

loop forever

q1: non-critical section

q2: wantq \leftarrow true

q3: while wantp

q4: if turn = 1

q5: wantq \leftarrow false

q6: await turn = 2

q7: wantq \leftarrow true

q8: critical section

q9: turn \leftarrow 1

q10: wantq \leftarrow false

Dekker's Algorithm

- Similar to fourth attempt
- *Right to insist on entering* rather than *Right to enter* is explicitly passed between the processes
- The individual variables ensure mutual exclusion
- Dekker's algorithm is correct: it satisfies the mutual exclusion property
- Its free from starvation and deadlock

Advanced Algorithm

- Two reasons
 - They work for arbitrary number of processes
 - They raise additional issues concerning concurrency abstraction and the specification of the CS problem
 - Developed by Leslie Lamport
- Bakery Algorithm

Bakery Algorithm

- A process wishing to enter its critical section is required to take a numbered ticket, whose value is greater than the values of all outstanding tickets
- The process waits until its ticket has the lowest value of all outstanding tickets and enters its critical section
- Name is taken from the ticket dispensers used in bakeries
- Customers are served first come first served basis

Bakery Algorithm

Algorithm 5.1: Bakery algorithm (two processes)

integer $np \leftarrow 0$, $nq \leftarrow 0$

p

loop forever

p1: non-critical section
p2: $np \leftarrow nq + 1$
p3: await $nq = 0$ or $np \leq nq$
p4: critical section
p5: $np \leftarrow 0$

q

loop forever

q1: non-critical section
q2: $nq \leftarrow np + 1$
q3: await $np = 0$ or $nq < np$
q4: critical section
q5: $nq \leftarrow 0$

Bakery Algorithm

- np and nq hold the ticket numbers of the two processes
- A value of 0 indicates that the process does not want to enter the CS
- A positive value represents an implicit queue of the processes that want to enter
- Lower numbers denoting closeness to the head of the queue
- If ticket numbers are equal we arbitrarily assign precedence to one process

Bakery Algorithm

- The following formulas are invariant
 - $np = 0 \leftrightarrow p1 \vee p2$
 - $nq = 0 \leftrightarrow q1 \vee q2$
 - $p4 \rightarrow (nq = 0) \vee (np \leq nq)$
 - $q4 \rightarrow (np = 0) \vee (nq < np)$
- Bakery Algorithm for two processes satisfies mutual exclusion property
 - $(p4 \wedge q4) \rightarrow (np \leq nq) \wedge (nq < np)$
- The two process bakery algorithm is free from starvation

Bakery Algorithm (N Processes)

Algorithm 5.2: Bakery algorithm (N processes)

integer array[1..n] number \leftarrow [0,...,0]

loop forever

p1: non-critical section

p2: number[i] \leftarrow 1 + max(number)

p3: for all *other* processes j

p4: await (number[j] = 0) or (number[i] \ll number[j])

p5: critical section

p6: number[i] \leftarrow 0

Bakery Algorithm (N Processes)

- Each of the n processes executes the same algorithm
- i is set to a different constant in the range 1 to N called the ID number of the processes
- The statement for all other processes j is an abbreviation for

```
for j from 1 to N  
    if j  $\neq$  i
```


Bakery Algorithm (N Processes)

- The notation $(\text{number}[i] \ll \text{number}[j])$ is an abbreviation for
$$(\text{number}[i] < \text{number}[j]) \text{ or } ((\text{number}[i] = \text{number}[j]) \text{ and } (i < j))$$
- Either the first ticket is less than the second or they are equal and first process ID number is lower than the second

Bakery Algorithm (N Processes)

- Each process chooses a number that is greater than the maximum of all outstanding ticket numbers
- A process is allowed to enter its critical section when it has lower ticket number than all other processes
- In case of tie in comparing the ticket numbers, the lower numbered process is arbitrarily given precedence.

Bakery Algorithm (N Processes)

- Algorithm is elegant because it has the property that no variable is both read and written by more than one process (unlike *turn* variable in Dekker's algorithm).
- However it is impractical because each process must query every other process for the value of its ticket number and the entire loop must be extended even if no other process wants to enter the critical section.