

POSIX THREADS PROGRAMMING

Dr. Lavika Goel

What is a Thread?

- Thread is defined as an independent stream of instructions that can be scheduled to run as such by the operating system
- Its a "procedure" that runs independently from its main program
- Multithreaded Program
 - Imagine a main program that contains a number of procedures
 - Imagine all of these procedures being able to be scheduled to run simultaneously or independently

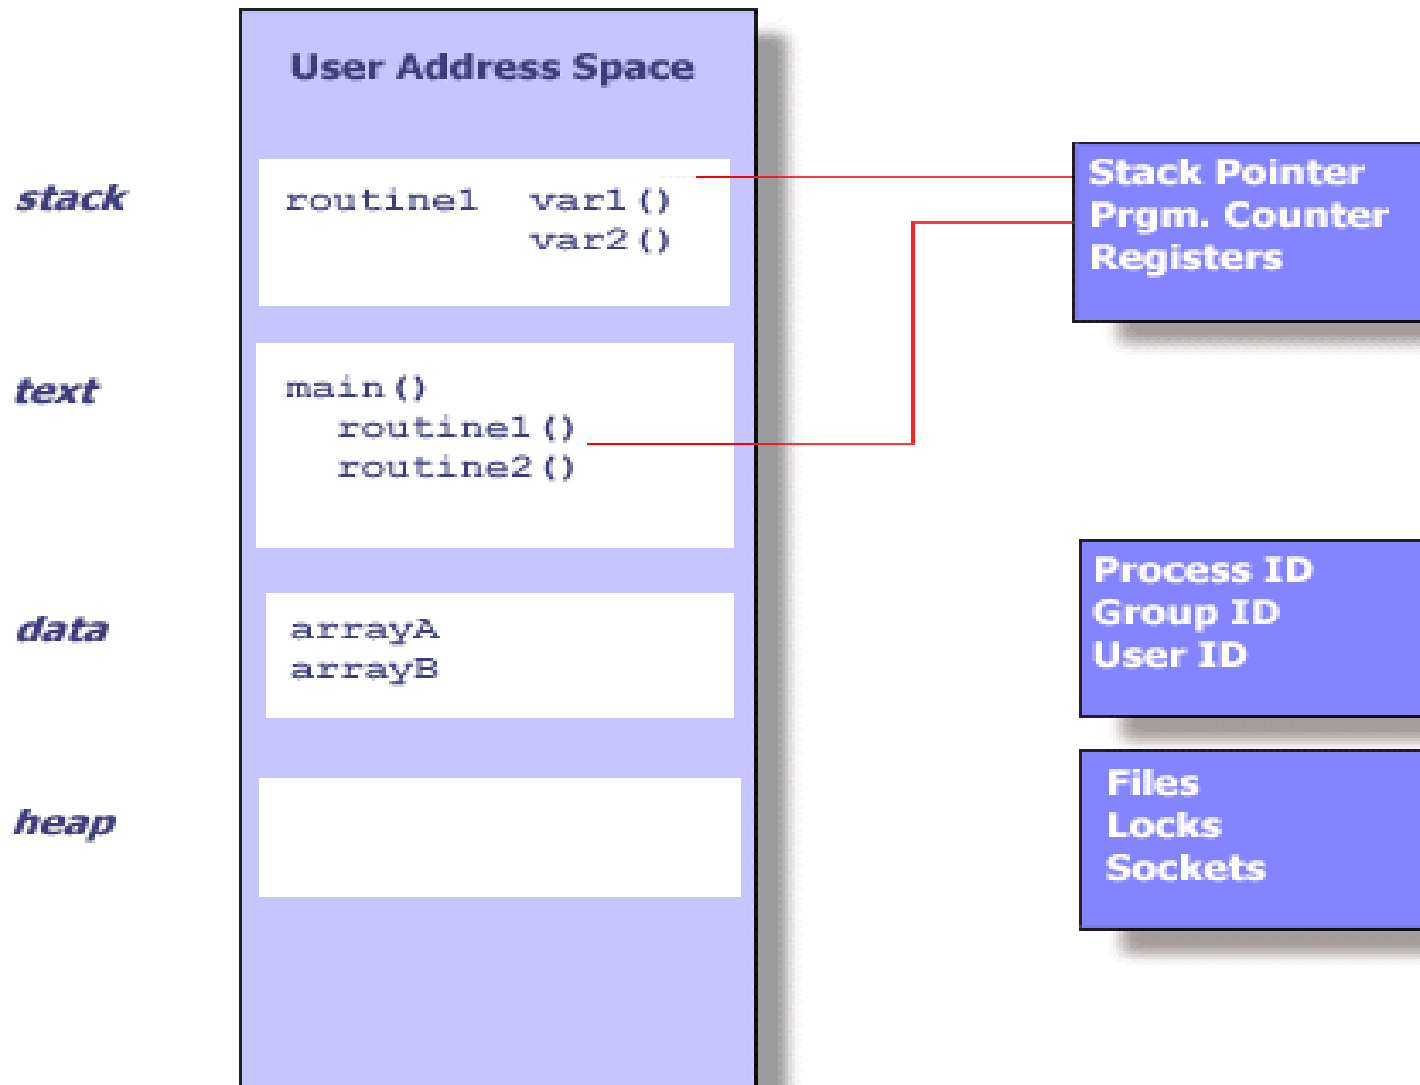
Process

- A process is created by the operating system, and requires a fair amount of "overhead".
- Processes contain information about program resources and program execution state
- Process Information:
 - Process ID, process group ID, user ID, and group ID
 - Environment
 - Working directory
 - Program instructions

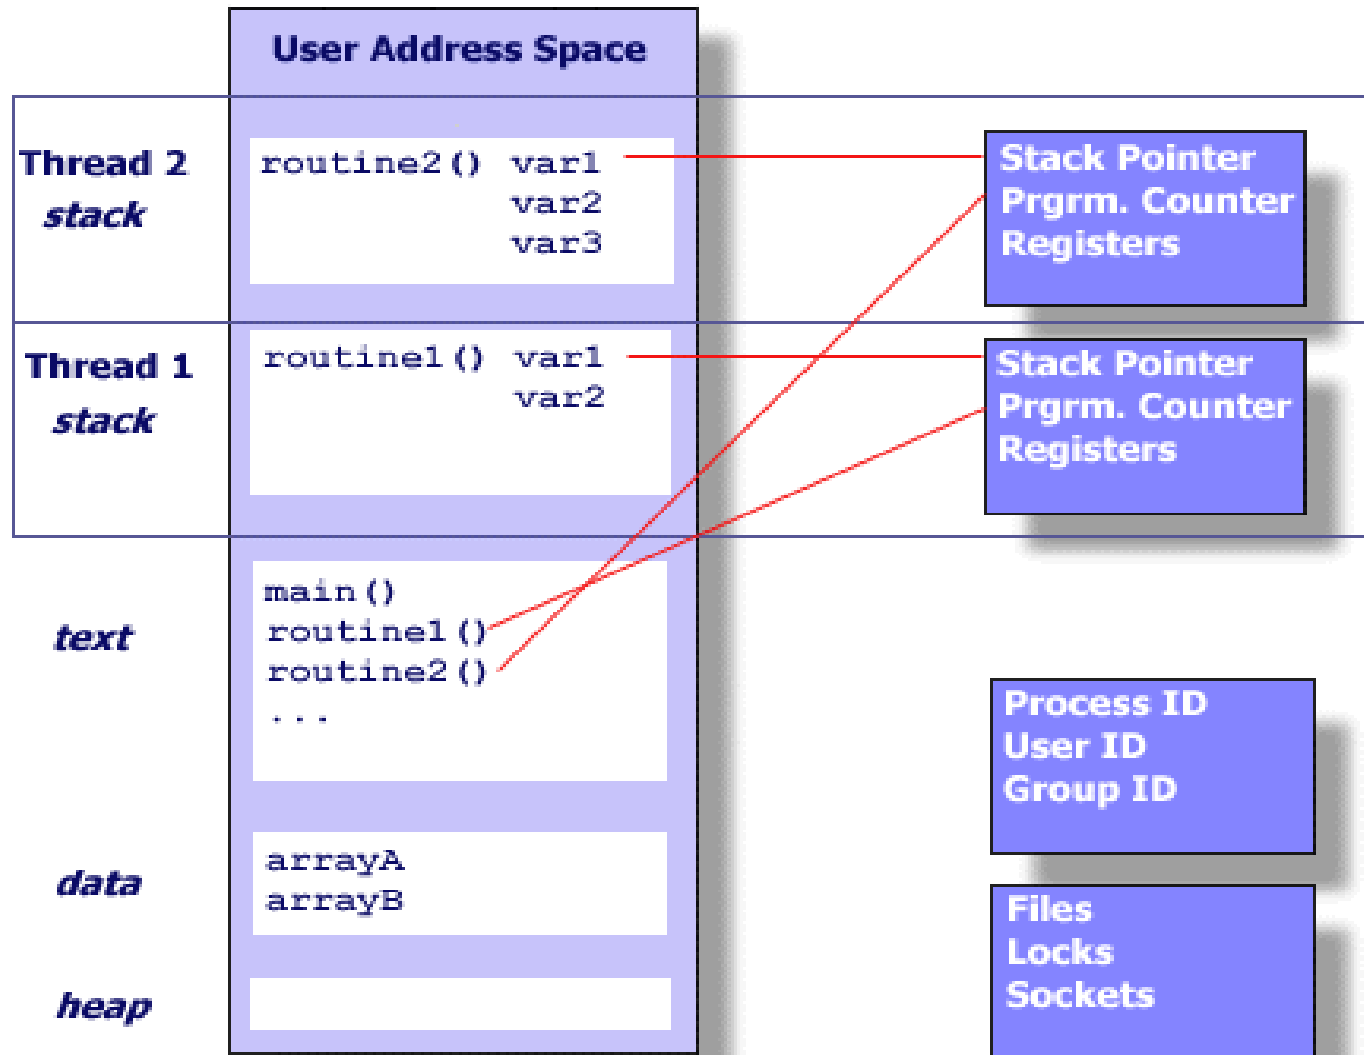
Process

- Process Information:
 - Registers, Stack, Heap
 - File descriptors
 - Signal actions
 - Shared libraries
 - Inter-process communication tools
 - message queues, pipes,
 - semaphores, or shared memory

Unix Process



Threads in a Unix Process



Thread Control

- Threads use and exist within the process resources
- Threads are able to be scheduled by the OS and run as independent entities
- Threads duplicate only the bare essential resources that enable them to exist as executable code
- Thread maintains its own Stack pointer, Registers, Scheduling properties, Set of pending and blocked signals and Thread specific data.

Threads in UNIX

- Has its own independent flow of control as long as its parent process exists and the OS supports it
- Duplicates only the essential resources it needs to be independently schedulable
- May share the process resources with other threads that act equally independently (and dependently)
- Dies if the parent process dies - or something similar
- Is "lightweight" because most of the overhead has already been accomplished through the creation of its process.

What are Pthreads?

- Hardware vendors have implemented their own proprietary versions of threads.
- These implementations differed substantially from each other making portable threaded applications difficult to develop.
- Standardized programming interface was required to take full advantage of the capabilities provided by threads
- IEEE POSIX 1003.1c standard (1995) - Portable Operating System Interface (POSIX)

Pthreads

- Implementations adhering to this standard are referred to as POSIX threads, or Pthreads.
- Most hardware vendors now offer Pthreads in addition to their proprietary API's.
- Pthreads are defined as a set of C language programming types and procedure calls, implemented with a **pthread.h** header / include file and a thread library
- This library may be part of another library, such as **libc**, in some implementations.

Some useful links:

- standards.ieee.org/findstds/standard/1003.1-2008.html
- www.opengroup.org/austin/papers/posix_faq.html
- www.unix.org/version3/ieee_std.html

Lightweight

- When compared to the cost of creating and managing a process, a thread can be created with much less operating system overhead.
- Managing threads requires fewer system resources than managing processes.
- For example, the following table compares timing results for the **fork()** subroutine and the **pthread_create()** subroutine.
- Timings reflect 50,000 process/thread creations, were performed with the **time** utility, and units are in seconds, no optimization flags.

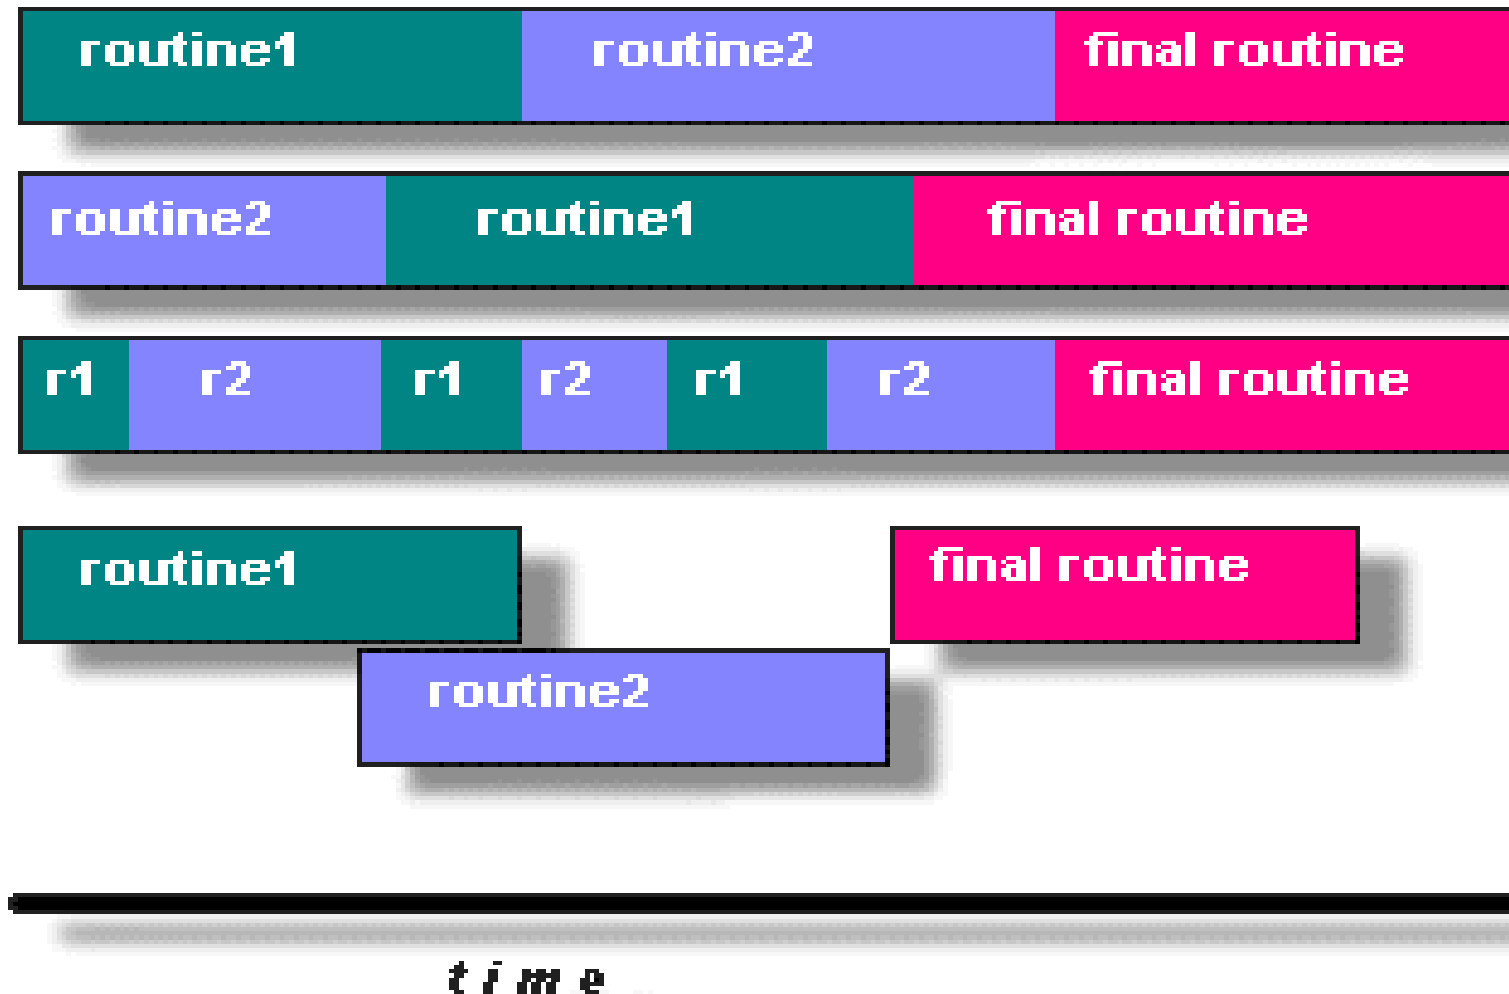
Lightweight

Platform	fork()			pthread_create()		
	real	user	sys	real	user	sys
Intel 2.6 GHz Xeon E5-2670 (16 cores/node)	8.1	0.1	2.9	0.9	0.2	0.3
Intel 2.8 GHz Xeon 5660 (12 cores/node)	4.4	0.4	4.3	0.7	0.2	0.5
AMD 2.3 GHz Opteron (16 cores/node)	12.5	1.0	12.5	1.2	0.2	1.3
AMD 2.4 GHz Opteron (8 cores/node)	17.6	2.2	15.7	1.4	0.3	1.3
IBM 4.0 GHz POWER6 (8 cpus/node)	9.5	0.6	8.8	1.6	0.1	0.4
IBM 1.9 GHz POWER5 p5-575 (8 cpus/node)	64.2	30.7	27.6	1.7	0.6	1.1
IBM 1.5 GHz POWER4 (8 cpus/node)	104.5	48.6	47.2	2.1	1.0	1.5
INTEL 2.4 GHz Xeon (2 cpus/node)	54.9	1.5	20.8	1.6	0.7	0.9
INTEL 1.4 GHz Itanium2 (4 cpus/node)	54.5	1.1	22.2	2.0	1.2	0.6

Concurrent and Parallel Programming

- Problem partitioning
- Load balancing
- Communications
- Data dependencies
- Synchronization and race conditions
- Memory issues
- I/O issues
- Program complexity
- Programmer effort/costs/time

Concurrent (Parallel) Execution

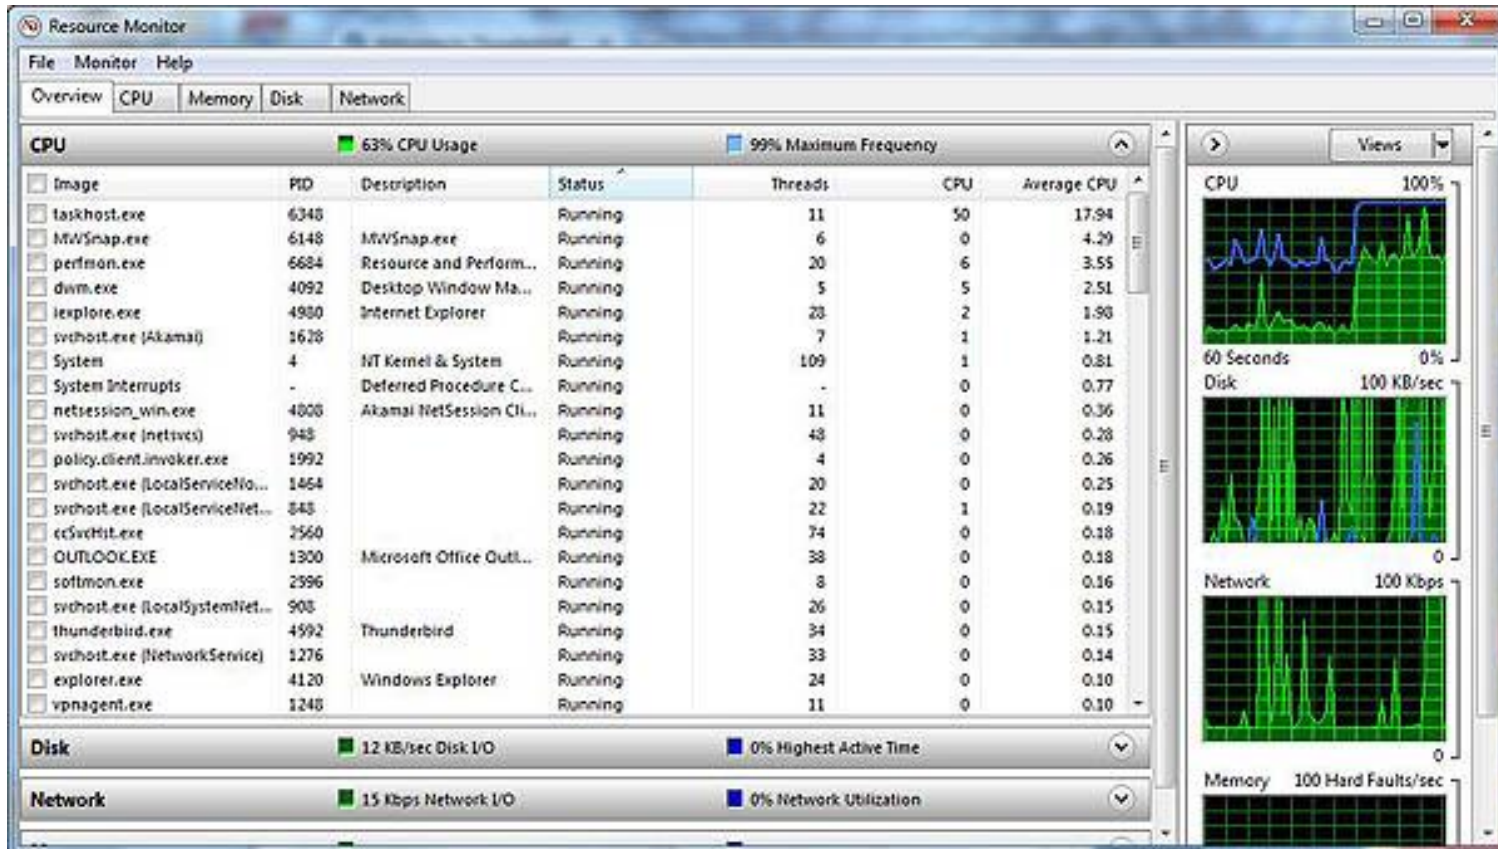


Programs well suited for pthreads


- Work that can be executed, or data that can be operated on, by multiple tasks simultaneously.
- Block for potentially long I/O waits
- Use many CPU cycles in some places but not others
- Must respond to asynchronous events
- Some work is more important than other work (priority interrupts)

Example

- A perfect example is the typical web browser, where many interleaved tasks can be happening at the same time, and where tasks can vary in priority.
 - Another good example is a modern operating system, which makes extensive use of threads. A screenshot of the MS Windows OS and applications using threads is shown below.



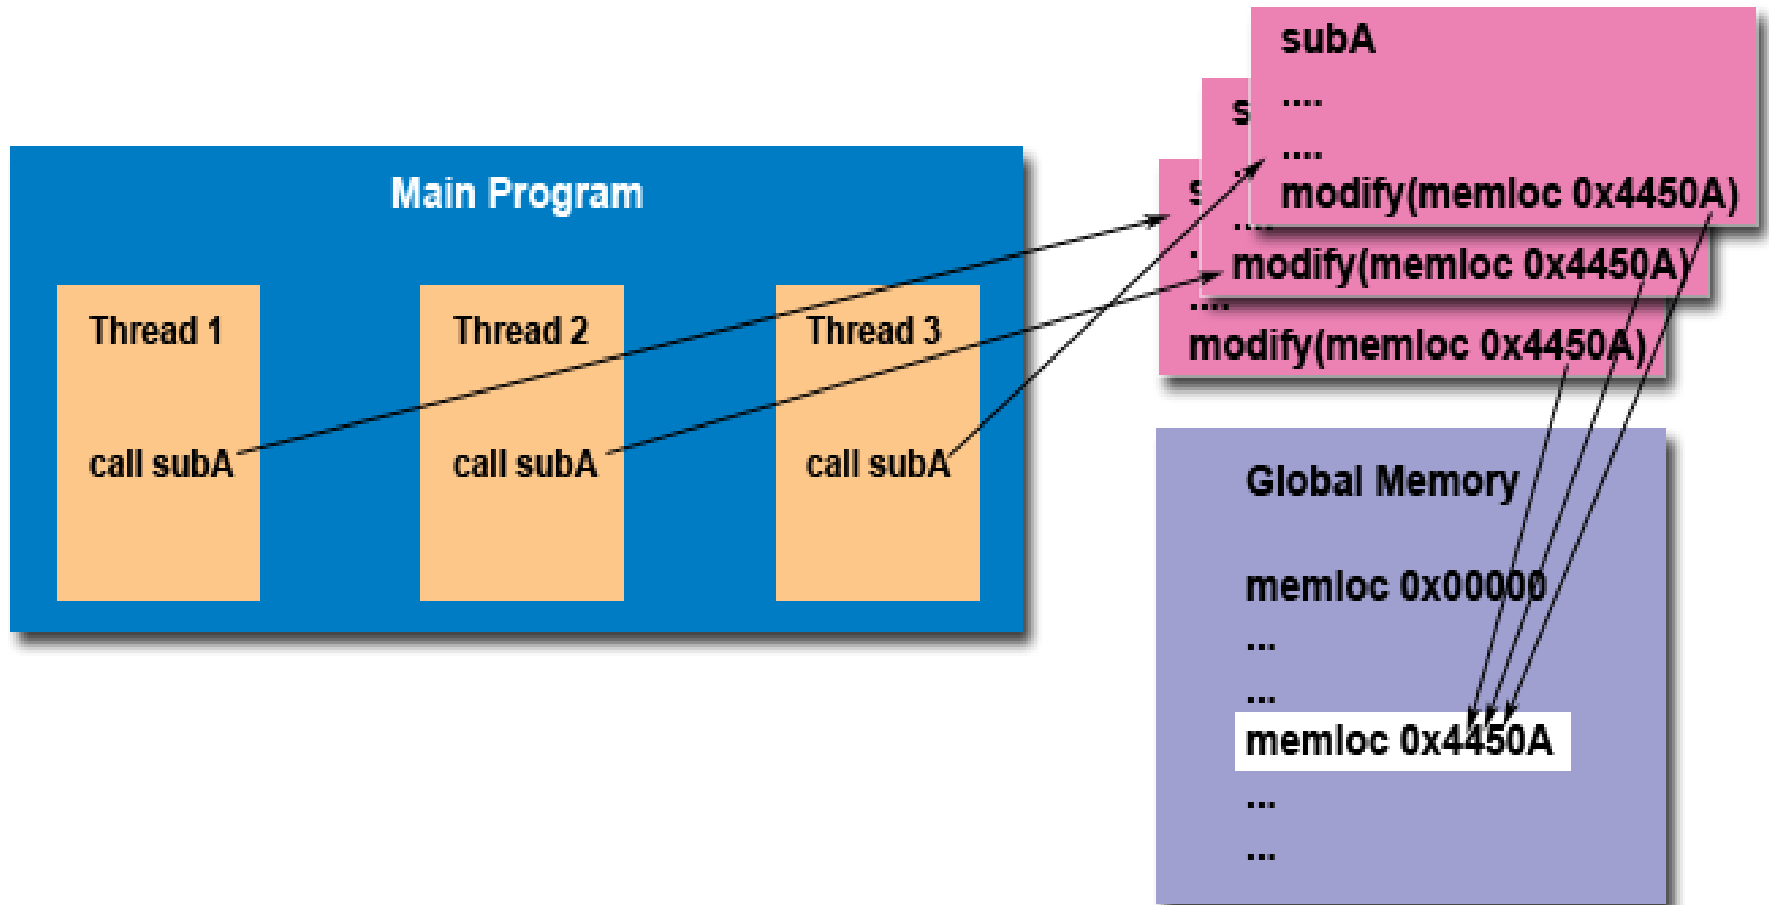
Shared Memory



Shared Memory Model

- All threads have access to the same global, shared memory
- Threads also have their own private data
- Programmers are responsible for synchronizing access (protecting) globally shared data.

Thread Safeness



Thread Safeness

- Thread-safeness refers to an application's ability to execute multiple threads simultaneously without "clobbering" shared data or creating "race" conditions
- If an application creates several threads, each of which makes a call to the same library routine:
 - This library routine modifies a global memory.
 - Each thread may try to modify this global memory location at the same time.
 - If the routine does not employ some sort of synchronization constructs, it is not thread-safe.

Pthreads API

- Pthreads API was defined in the ANSI/IEEE POSIX 1003.1 - 1995
- Four major Categories
 - **Thread management:** Creating, Detaching, Joining
 - **Mutexes** (mutual exclusion) deals with synchronization: Mutex creation, destroying, locking, unlocking
 - **Condition variables**
 - **Synchronization:** Routines that manage read / write locks and barriers

Pthread Library

Routine Prefix	Functional Group
pthread_	Threads themselves and miscellaneous subroutines
pthread_attr_	Thread attributes objects
pthread_mutex_	Mutexes
pthread_mutexattr_	Mutex attributes objects.
pthread_cond_	Condition variables
pthread_condattr_	Condition attributes objects
pthread_key_	Thread-specific data keys
pthread_rwlock_	Read/write locks
pthread_barrier_	Synchronization barriers

Pthread Library

- The Pthreads API contains around 100 subroutines
- For portability, the **pthread.h** header file should be included in each source file using the Pthreads library.
- The current POSIX standard is defined only for the C language.
- Compile Commands - GNU C
gcc -pthread

Creating and Managing Threads

- `pthread_create (thread, attr, start_routine, arg)`
- `pthread_exit (status)`
- `pthread_cancel (thread)`
- `pthread_attr_init (attr)`
- `pthread_attr_destroy (attr)`


Creating Threads

- `main()` program comprises a single, default thread
- All other threads must be explicitly created by the programmer
- `pthread_create` creates a new thread and makes it executable
- This routine can be called any number of times from anywhere within your code
- Once created, threads are peers, and may create other threads
- There is no implied hierarchy or dependency between threads.

pthread_create – Arguments

- `thread`: A unique identifier for the new thread returned by the subroutine.
- `attr`: An attribute object that may be used to set thread attributes. We can specify a thread attributes object, or NULL for the default values.
- `start_routine`: C routine that the thread will execute once it is created.
- `arg`: A single argument that may be passed to *start_routine*. It must be passed by reference as a pointer cast of type void. NULL may be used if no argument is to be passed

Creating Threads



Thread creation and termination example

This simple example code creates 5 threads with the `pthread_create()` routine. Each thread prints a "Hello World!" message, and then terminates with a call to `pthread_exit()`.



Pthread Creation and Termination Example

```
1  #include <pthread.h>
2  #include <stdio.h>
3  #define NUM_THREADS      5
4
5  void *PrintHello(void *threadid)
6  {
7      long tid;
8      tid = (long)threadid;
9      printf("Hello World! It's me, thread #%ld!\n",
10 tid);
11     pthread_exit(NULL);
12 }
13
14 int main (int argc, char *argv[])
15 {
16     pthread_t threads[NUM_THREADS];
17     int rc;
18     long t;
19     for(t=0; t<NUM_THREADS; t++){
20         printf("In main: creating thread %ld\n", t);
21         rc = pthread_create(&threads[t], NULL,
22 PrintHello, (void *)t);
23         if (rc){
24             printf("ERROR; return code from
25 pthread_create() is %d\n", rc);
26             exit(-1);
27         }
28     }
29
30     /* Last thing that main() should do */
31     pthread_exit(NULL);
32 }
```

Thread Attributes

- Thread is created with certain default attributes. Some of these attributes can be changed by the programmer via the thread attribute object.
- `pthread_attr_init` & `pthread_attr_destroy` are used to initialize / destroy the thread attribute object.
- Other routines are then used to query / set specific attributes in the thread attribute object
- Attributes include:
 - Detached or joinable state
 - Scheduling inheritance, policy, parameters, contention scope
 - Stack size, address, guard (overflow) size

Terminating Threads

- The thread returns normally from its starting routine. Its work is done.
- The thread makes a call to the `pthread_exit` subroutine - whether its work is done or not.
- The thread is canceled by another thread via the `pthread_cancel` routine.
- The entire process is terminated due to making a call to either the `exec()` or `exit()`
- If `main()` finishes first, without calling `pthread_exit` explicitly itself

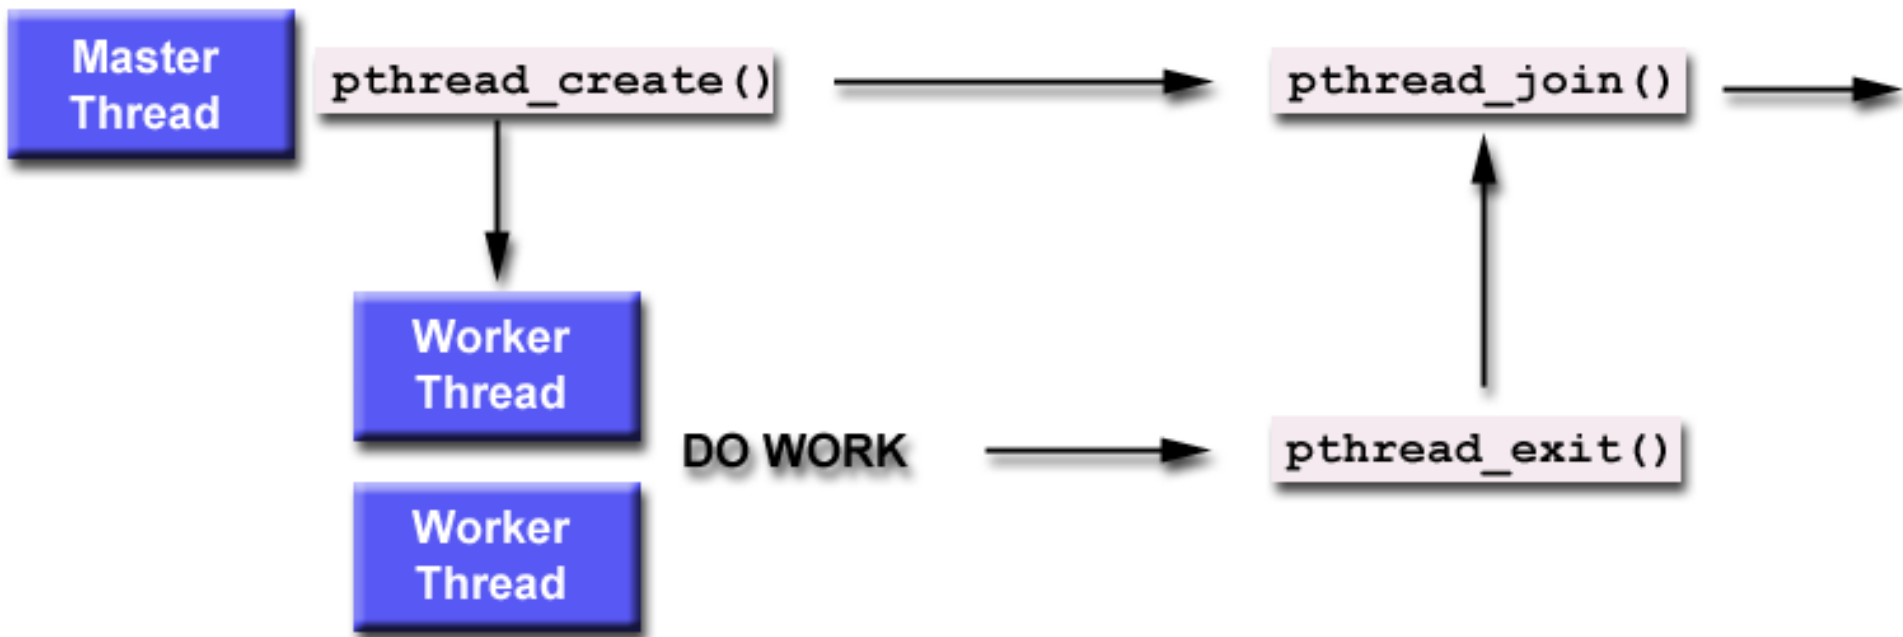
Terminating Threads

- The `pthread_exit()` routine allows the programmer to specify an optional termination *status* parameter
- This optional parameter is typically returned to threads "joining" the terminated thread
- In subroutines that execute to completion normally, calling `pthread_exit()` can often be dispensed with unless the optional status code needs to be passed back
- Cleanup: the `pthread_exit()` routine does not close files; any files opened inside the thread will remain open after the thread is terminated

Joining and Detaching Threads

- `pthread_join (threadid, status)`
- `pthread_detach (threadid)`
- `pthread_attr_setdetachstate (attr, detachstate)`
- `pthread_attr_getdetachstate (attr, detachstate)`

Joining



Joining

- "Joining" is one way to accomplish synchronization between threads
- The `pthread_join()` subroutine blocks the calling thread until the specified threadid thread terminates.
- The programmer is able to obtain the target thread's termination return status if it was specified in the target thread's call to `pthread_exit()`.
- A joining thread can match one `pthread_join()` call
- It is a logical error to attempt multiple joins on the same thread.

Joinable?

- When a thread is created, one of its attributes defines whether it is joinable or detached
- Only threads that are created as joinable can be joined
- If a thread is created as detached, it can never be joined
- The final draft of the POSIX standard specifies that threads should be created as joinable
- To explicitly create a thread as joinable or detached, the `attr` argument in the `pthread_create()` routine is used

Joinable?

- Four Step Process:
 - Declare a pthread attribute variable of the pthread_attr_t data type
 - Initialize the attribute variable with **pthread_attr_init()**
 - Set the attribute detached status with **pthread_attr_setdetachstate()**
 - When done, free library resources used by the attribute with **pthread_attr_destroy()**

Detaching

- The `pthread_detach()` routine can be used to explicitly detach a thread even though it was created as joinable
- There is no converse routine

Example: Pthread Joining

- This example demonstrates how to "wait" for thread completions by using the Pthread join routine.
- Since some implementations of Pthreads may not create threads in a joinable state, the threads in this example are explicitly created in a joinable state so that they can be joined later.

Example



Pthread Joining Example

```
1  #include <pthread.h>
2  #include <stdio.h>
3  #include <stdlib.h>
4  #include <math.h>
5  #define NUM_THREADS      4
6
7  void *BusyWork(void *t)
8  {
9      int i;
10     long tid;
11     double result=0.0;
12     tid = (long)t;
13     printf("Thread %ld starting...\n",tid);
14     for (i=0; i<1000000; i++)
15     {
16         result = result + sin(i) * tan(i);
17     }
18     printf("Thread %ld done. Result = %e\n",tid,
19 result);
```


(Contd.)

```
20     pthread_exit((void*) t);
21 }
22
23 int main (int argc, char *argv[])
24 {
25     pthread_t thread[NUM_THREADS];
26     pthread_attr_t attr;
27     int rc;
28     long t;
29     void *status;
30
31     /* Initialize and set thread detached attribute */
32     pthread_attr_init(&attr);
33     pthread_attr_setdetachstate(&attr,
34 PTHREAD_CREATE_JOINABLE);
35
36     for(t=0; t<NUM_THREADS; t++) {
37         printf("Main: creating thread %ld\n", t);
38         rc = pthread_create(&thread[t], &attr,
39 BusyWork, (void *)t);
40         if (rc) {
41             printf("ERROR; return code from
42 pthread_create() is %d\n", rc);
43             exit(-1);
44         }
45     }
46
47     /* Free attribute and wait for the other threads
48 */
49     pthread_attr_destroy(&attr);
50     for(t=0; t<NUM_THREADS; t++) {
51         rc = pthread_join(thread[t], &status);
52         if (rc) {
53             printf("ERROR; return code from
54 pthread_join() is %d\n", rc);
55             exit(-1);
56         }
57         printf("Main: completed join with thread %ld
having a status
of %ld\n", t, (long)status);
    }

    printf("Main: program completed. Exiting.\n");
    pthread_exit(NULL);
}
```

Mutex Variables

- Mutex is an abbreviation for "mutual exclusion".
- Mutex variables are one of the primary means of implementing thread synchronization and for protecting shared data when multiple writes occur.
- A mutex variable acts like a "lock" protecting access to a shared data resource.

Mutex Variables

- The basic concept of a mutex as used in Pthreads is that only one thread can lock (or own) a mutex variable at any given time.
- Even if several threads try to lock a mutex only one thread will be successful.
- No other thread can own that mutex until the owning thread unlocks that mutex.
- Threads must "take turns" accessing protected data.

Race Conditions

- Mutexes can be used to prevent "race" conditions.

Thread 1	Thread 2	Balance
Read balance: \$1000		\$1000
	Read balance: \$1000	\$1000
	Deposit \$200	\$1000
Deposit \$200		\$1000
Update balance $\$1000 + \200		\$1200
	Update balance $\$1000 + \200	\$1200

Mutex

- Create and initialize a mutex variable
- Several threads attempt to lock the mutex
- Only one succeeds and that thread owns the mutex
- The owner thread performs some set of actions
- The owner unlocks the mutex
- Another thread acquires the mutex and repeats the process
- Finally the mutex is destroyed

Create and Destroy Mutex

- `pthread_mutex_init (mutex, attr)`
- `pthread_mutex_destroy (mutex)`
- `pthread_mutexattr_init (attr)`
- `pthread_mutexattr_destroy (attr)`

Create and Destroy Mutex

- Mutex variables must be declared with type `pthread_mutex_t`, and must be initialized before they can be used.
- There are two ways to initialize a mutex variable:
 - Statically, when it is declared

```
pthread_mutex_t mymutex =  
PTHREAD_MUTEX_INITIALIZER;
```
 - Dynamically, with the `pthread_mutex_init()` routine. This method permits setting mutex object attributes, *attr*
- The mutex is initially unlocked.

Create and Destroy Mutex

- The *attr* object is used to establish properties for the mutex object, and must be of type `pthread_mutexattr_t` if used (may be specified as NULL to accept defaults).
- The `pthread_mutexattr_init()` and `pthread_mutexattr_destroy()` routines are used to create and destroy mutex attribute objects respectively.
- `pthread_mutex_destroy()` should be used to free a mutex object which is no longer needed.

Lock and Unlock Mutex

- `pthread_mutex_lock (mutex)`
-
- `pthread_mutex_trylock (mutex)`
- `pthread_mutex_unlock (mutex)`

Lock and Unlock Mutex

- The `pthread_mutex_lock()` routine is used by a thread to acquire a lock on the *mutex* variable.
- If the mutex is already locked by another thread, this call will block the calling thread until the mutex is unlocked.
- `pthread_mutex_trylock()` will lock a mutex.
- However, if the mutex is already locked, the routine will return immediately with a "busy" error code.
- This routine may be useful in preventing deadlock conditions, as in a priority-inversion situation.

Lock and Unlock Mutex

- `pthread_mutex_unlock()` will unlock a mutex if called by the owning thread.
- Calling this routine is required after a thread has completed its use of protected data and other threads need access.
- An error will be returned if:
 - If the mutex was already unlocked
 - If the mutex is owned by another thread
- Mutexes are akin to a "gentlemen's agreement"
- It is up to the code writer to insure that the necessary threads all make the mutex lock and unlock calls correctly

Example: Using Mutexes

- This example program illustrates the use of mutex variables in a threads program that performs a dot product.
- The main data is made available to all threads through a globally accessible structure.
- Each thread works on a different part of the data.
- The main thread waits for all the threads to complete their computations, and then it prints the resulting sum.

Example



Using Mutexes Example

```
1 #include <pthread.h>
2 #include <stdio.h>
3 #include <stdlib.h>
4
5 /*
6  The following structure contains the necessary
7 information
8 to allow the function "dotprod" to access its input
9 data and
10 place its output into the structure.
11 */
12
13 typedef struct
14 {
15     double    *a;
16     double    *b;
17     double    sum;
18     int       veclen;
19 } DOTDATA;
20
21 /* Define globally accessible variables and a mutex
22 */
23
24 #define NUMTHRDS 4
25 #define VECLEN 100
26     DOTDATA dotstr;
27     pthread_t callThd[NUMTHRDS];
28     pthread_mutex_t mutexsum;
29
30 /*
31 The function dotprod is activated when the thread is
32 created.
33 All input to this routine is obtained from a
34 structure
35 of type DOTDATA and all output from this function is
36 written into
37 this structure. The benefit of this approach is
38 apparent for the
39 multi-threaded program: when a thread is created we
40 pass a single
41 argument to the activated function - typically this
42 argument
43 is a thread number. All the other information
```

(Contd. - 1)

```
44 required by the
45 function is accessed from the globally accessible
46 structure.
47 */
48
49 void *dotprod(void *arg)
50 {
51
52     /* Define and use local variables for convenience
53 */
54
55     int i, start, end, len ;
56     long offset;
57     double mysum, *x, *y;
58     offset = (long)arg;
59
60     len = dotstr.vecilen;
61     start = offset*len;
62     end   = start + len;
63     x = dotstr.a;
64     y = dotstr.b;
65
66     /*
67     Perform the dot product and assign result
68     to the appropriate variable in the structure.
69     */
70
71     mysum = 0;
72     for (i=start; i<end ; i++)
73     {
74         mysum += (x[i] * y[i]);
75     }
76
77     /*
78     Lock a mutex prior to updating the value in the
79 shared
80 structure, and unlock it upon updating.
81     */
82     pthread_mutex_lock (&mutexsum);
83     dotstr.sum += mysum;
84     pthread_mutex_unlock (&mutexsum);
85
86     pthread_exit((void*) 0);
87 }
88
```

(Contd. – 2)

```
89  /*
90  The main program creates threads which do all the
91  work and then
92  print out result upon completion. Before creating the
93  threads,
94  the input data is created. Since all threads update a
95  shared structure,
96  we need a mutex for mutual exclusion. The main thread
97  needs to wait for
98  all threads to complete, it waits for each one of the
99  threads. We specify
100  a thread attribute value that allow the main thread
101  to join with the
102  threads it creates. Note also that we free up handles
103  when they are
104  no longer needed.
105  */
106
107  int main (int argc, char *argv[])
108  {
109      long i;
110      double *a, *b;
111      void *status;
112      pthread_attr_t attr;
113
114      /* Assign storage and initialize values */
115      a = (double*) malloc
116  (NUMTHRDS*VECLEN*sizeof(double));
117      b = (double*) malloc
118  (NUMTHRDS*VECLEN*sizeof(double));
119
120      for (i=0; i<VECLEN*NUMTHRDS; i++)
121      {
122          a[i]=1.0;
123          b[i]=a[i];
124      }
125
126      dotstr.vecLEN = VECLen;
127      dotstr.a = a;
128      dotstr.b = b;
129      dotstr.sum=0;
130
131      pthread_mutex_init(&mutexsum, NULL);
132
133      /* Create threads to perform the dotproduct */
```

(Contd. – 3)

```
134     pthread_attr_init(&attr);
135     pthread_attr_setdetachstate(&attr,
136 PTHREAD_CREATE_JOINABLE);
137
138     for(i=0; i<NUMTHRDS; i++)
    {
        /*
        Each thread works on a different set of data. The
offset is specified
        by 'i'. The size of the data for each thread is
indicated by VECLEN.
        */
        pthread_create(&callThd[i], &attr, dotprod, (void
*)i);
    }

    pthread_attr_destroy(&attr);

    /* Wait on the other threads */
    for(i=0; i<NUMTHRDS; i++)
    {
        pthread_join(callThd[i], &status);
    }

    /* After joining, print out the results and
cleanup */
    printf ("Sum =  %f \n", dotstr.sum);
    free (a);
    free (b);
    pthread_mutex_destroy(&mutexsum);
    pthread_exit(NULL);
}
```


Semaphore

- The pthreads library itself does not provide a semaphore
- However, a separate POSIX standard does define them
- The necessary declarations to use these semaphores are contained in semaphore.h
- To define a semaphore object, use
`sem_t sem_name;`

Semaphore API

- **sem_init**: Initialize a new semaphore
 - The second argument denotes *how* the semaphore will be shared.
 - Passing zero denotes that it will be shared among **threads** rather than processes.
 - The final argument is the initial value of the semaphore.
- **sem_destroy**: Deallocate an existing semaphore.
- **sem_wait**: This is the Wait () operation.
- **sem_post**: This is the Signal () operation.

sem_init

- `int sem_init (sem_t *sem, int pshared, unsigned int value);`
 - `sem` address of the declared semaphore
 - `pshared` should be 0 (not shared with threads in other processes)
 - `value` the desired initial value of the semaphore
- Return: The return value is 0 if successful.

sem_wait

- To wait on a semaphore
- `int sem_wait (sem_t *sem);`
- If the value of the semaphore is negative, the calling process blocks
- One of the blocked processes wakes up when another process calls `sem_post`

wait (S)

```
    if S.V > 0
        S.V ← S.V - 1
    else
        S.L ← append (S.L , p)
        p.state ← blocked
```

sem_post

- To increment the value of a semaphore
- `int sem_post(sem_t *sem);`
- It increments the value of the semaphore
- It wakes up a blocked process waiting on the semaphore, if any.

signal (S)

if $S.L = \phi$

$S.V \leftarrow S.V + 1$

else

Let q be some process in $S.L$

$S.L \leftarrow S.L - \{q\}$

$q.state \leftarrow \text{ready}$

sem_getvalue

- To find out the value of a semaphore
- `int sem_getvalue (sem_t *sem, int *valp) ;`
- Gets the current value of sem
- Places it in the location pointed to by valp

sem_destroy

- To destroy a semaphore
- `sem_destroy(sem_t *sem) ;`
- Destroys the semaphore
- No threads should be waiting on the semaphore if its destruction is to succeed.

Comparison

State	Pthread	Mutex	Semaphore
Creation	<code>pthread_create</code>	<code>pthread_mutex_init</code>	<code>sem_init</code>
Destroy	<code>pthread_exit</code>	<code>pthread_mutex_destroy</code>	<code>sem_destroy</code>
Waiting	<code>pthread_join</code>	-	-
Acquisition	-	<code>pthread_mutex_lock</code>	<code>sem_wait</code>
Release	-	<code>pthread_mutex_unlock</code>	<code>sem_post</code>

Example code: Dining Philosopher Problem Using Semaphores

```
#include <pthread.h>
#include <semaphore.h>
#include <stdio.h>

#define N 5
#define THINKING 2
#define HUNGRY 1
#define EATING 0
#define LEFT (phnum + 4) % N
#define RIGHT (phnum + 1) % N

int state[N];
int phil[N] = { 0, 1, 2, 3, 4 };

sem_t mutex;
sem_t S[N];
```

(Contd.-1)

```
void test(int phnum)
{
    if (state[phnum] == HUNGRY
        && state[LEFT] != EATING
        && state[RIGHT] != EATING) {
        // state that eating
        state[phnum] = EATING;

        sleep(2);

        printf("Philosopher %d takes fork %d and %d\n",
               phnum + 1, LEFT + 1, phnum + 1);

        printf("Philosopher %d is Eating\n", phnum + 1);

        // sem_post(&S[phnum]) has no effect
        // during takefork
        // used to wake up hungry philosophers
        // during putfork
        sem_post(&S[phnum]);
    }
}
```

(Contd.-2)

```
// take up chopsticks
void take_fork(int phnum)
{
    sem_wait(&mutex);

    // state that hungry
    state[phnum] = HUNGRY;

    printf("Philosopher %d is Hungry\n", phnum + 1);

    // eat if neighbours are not eating
    test(phnum);

    sem_post(&mutex);

    // if unable to eat wait to be signalled
    sem_wait(&S[phnum]);

    sleep(1);
}
```

(Contd.-3)

```
// put down chopsticks
void put_fork(int phnum)
{
    sem_wait(&mutex);

    // state that thinking
    state[phnum] = THINKING;

    printf("Philosopher %d putting fork %d and %d down\n",
           phnum + 1, LEFT + 1, phnum + 1);
    printf("Philosopher %d is thinking\n", phnum + 1);

    test(LEFT);
    test(RIGHT);

    sem_post(&mutex);
}
```

```
void* philosopher(void* num)
{
    while (1) {
        int* i = num;

        sleep(1);

        take_fork(*i);

        sleep(0);

        put_fork(*i);
    }
}
```

(Contd.- 4)

```
int main()
{
    int i;
    pthread_t thread_id[N];

    // initialize the semaphores
    sem_init(&mutex, 0, 1);

    for (i = 0; i < N; i++)
        sem_init(&S[i], 0, 0);

    for (i = 0; i < N; i++) {
        // create philosopher processes
        pthread_create(&thread_id[i], NULL,
                      philosopher, &phil[i]);

        printf("Philosopher %d is thinking\n", i + 1);
    }

    for (i = 0; i < N; i++)
        pthread_join(thread_id[i], NULL);
}
```

Condition Variables

- Condition variables provide yet another way for threads to synchronize.
- While mutexes implement synchronization by controlling thread access to data, condition variables allow threads to synchronize based upon the actual value of data.
- Without condition variables, the programmer would need to have threads continually polling (possibly in a critical section), to check if the condition is met. A condition variable is a way to achieve the same goal without polling.
- A condition variable is always used in conjunction with a mutex lock.

Creating and Destroying Condition Variables

- `pthread_cond_init (condition, attr)`
- `pthread_cond_destroy (condition)`
- `pthread_condattr_init (attr)`
- `pthread_condattr_destroy (attr)`

Creating and Destroying Condition Variables

- Condition variables must be declared with type **pthread_cond_t**, and must be initialized before they can be used. There are two ways to initialize a condition variable:
 - Statically, when it is declared. For example:
**pthread_cond_t myconvar =
PTHREAD_COND_INITIALIZER;**
 - Dynamically, with the **pthread_cond_init()** routine. The ID of the created condition variable is returned to the calling thread through the *condition* parameter.

Creating and Destroying Condition Variables

- The optional *attr* object is used to set condition variable attributes. There is only one attribute defined for condition variables: process-shared, which allows the condition variable to be seen by threads in other processes.
- The attribute object, if used, must be of type **pthread_condattr_t** (may be specified as NULL to accept defaults).
- The **pthread_condattr_init()** and **pthread_condattr_destroy()** routines are used to create and destroy condition variable attribute objects.
- **pthread_cond_destroy()** should be used to free a condition variable that is no longer needed.

Waiting and Signaling on Condition Variables

- `pthread_cond_wait (condition, mutex)`
- `pthread_cond_signal (condition)`
- `pthread_cond_broadcast (condition)`

Waiting and Signaling on Condition Variables

- **pthread_cond_wait()** blocks the calling thread until the specified *condition* is signalled. This routine should be called while *mutex* is locked, and it will automatically release the mutex while it waits. After signal is received and thread is awakened, *mutex* will be automatically locked for use by the thread.

Waiting and Signaling on Condition Variables

- The **pthread_cond_signal()** routine is used to signal (or wake up) another thread which is waiting on the condition variable. It should be called after *mutex* is locked, and must unlock *mutex* in order for **pthread_cond_wait()** routine to complete.
- The **pthread_cond_broadcast()** routine should be used instead of **pthread_cond_signal()** if more than one thread is in a blocking wait state.
- It is a logical error to call **pthread_cond_signal()** before calling **pthread_cond_wait()**.

Waiting and Signaling on Condition Variables

- Proper locking and unlocking of the associated mutex variable is essential when using these routines. For example:
- Failing to lock the mutex before calling **pthread_cond_wait()** may cause it NOT to block.
- Failing to unlock the mutex after calling **pthread_cond_signal()** may not allow a matching **pthread_cond_wait()** routine to complete (it will remain blocked).

A
representative
sequence for
using
condition
variables is
shown here:

Main Thread

- Declare and initialize global data/variables which require synchronization (such as "count")
- Declare and initialize a condition variable object
- Declare and initialize an associated mutex
- Create threads A and B to do work

Thread A

- Do work up to the point where a certain condition must occur (such as "count" must reach a specified value)
- Lock associated mutex and check value of a global variable
- Call `pthread_cond_wait()` to perform a blocking wait for signal from Thread-B. Note that a call to `pthread_cond_wait()` automatically and atomically unlocks the associated mutex variable so that it can be used by Thread-B.
- When signalled, wake up. Mutex is automatically and atomically locked.
- Explicitly unlock mutex
- Continue

Thread B

- Do work
- Lock associated mutex
- Change the value of the global variable that Thread-A is waiting upon.
- Check value of the global Thread-A wait variable. If it fulfills the desired condition, signal Thread-A.
- Unlock mutex.
- Continue

Main Thread

Join / Continue

Example: Using Condition Variables

- This simple example code demonstrates the use of several Pthread condition variable routines.
- The main routine creates three threads.
- Two of the threads perform work and update a "count" variable.
- The third thread waits until the count variable reaches a specified value.

Example



Using Condition Variables Example

```
1 #include <pthread.h>
2 #include <stdio.h>
3 #include <stdlib.h>
4
5 #define NUM_THREADS 3
6 #define TCOUNT 10
7 #define COUNT_LIMIT 12
8
9 int count = 0;
10 int thread_ids[3] = {0,1,2};
11 pthread_mutex_t count_mutex;
12 pthread_cond_t count_threshold_cv;
13
14 void *inc_count(void *t)
15 {
16     int i;
17     long my_id = (long)t;
18
19     for (i=0; i<TCOUNT; i++) {
20         pthread_mutex_lock(&count_mutex);
21         count++;
22
23         /*
24          Check the value of count and signal waiting
25 thread when condition is
26 reached. Note that this occurs while mutex is
27 locked.
28 */
```


(Contd.)

```
29     if (count == COUNT_LIMIT) {
30         pthread_cond_signal(&count_threshold_cv);
31         printf("inc_count(): thread %ld, count = %d\n",
32             my_id, count);
33     }
34     printf("inc_count(): thread %ld, count = %d,
35     unlocking mutex\n",
36         my_id, count);
37     pthread_mutex_unlock(&count_mutex);
38
39     /* Do some "work" so threads can alternate on
40     mutex lock */
41     sleep(1);
42 }
43 pthread_exit(NULL);
44 }
45
46 void *watch_count(void *t)
47 {
48     long my_id = (long)t;
49
50     printf("Starting watch_count(): thread %ld\n",
51         my_id);
52
53     /*
54     Lock mutex and wait for signal. Note that the
55     pthread_cond_wait
56     routine will automatically and atomically unlock
57     mutex while it waits.
58     Also, note that if COUNT_LIMIT is reached before
59     this routine is run by
60     the waiting thread, the loop will be skipped to
61     prevent pthread_cond_wait
62     from never returning.
63     */
64     pthread_mutex_lock(&count_mutex);
65     while (count < COUNT_LIMIT) {
66         pthread_cond_wait(&count_threshold_cv,
67             &count_mutex);
68         printf("watch_count(): thread %ld Condition
69         signal received.\n", my_id);
70         count += 125;
71         printf("watch_count(): thread %ld count now =
72         %d.\n", my_id, count);
73     }
```

(Contd.)

```
74     }
75     pthread_mutex_unlock(&count_mutex);
76     pthread_exit(NULL);
77 }
78
79 int main (int argc, char *argv[])
80 {
81     int i, rc;
82     long t1=1, t2=2, t3=3;
83     pthread_t threads[3];
84     pthread_attr_t attr;
85
86     /* Initialize mutex and condition variable objects
87 */
88     pthread_mutex_init(&count_mutex, NULL);
89     pthread_cond_init (&count_threshold_cv, NULL);
90
91     /* For portability, explicitly create threads in a
92 joinable state */
93     pthread_attr_init(&attr);
94     pthread_attr_setdetachstate(&attr,
95 PTHREAD_CREATE_JOINABLE);
96     pthread_create(&threads[0], &attr, watch_count,
97 (void *)t1);
98     pthread_create(&threads[1], &attr, inc_count, (void
99 *)t2);
100    pthread_create(&threads[2], &attr, inc_count, (void
101 *)t3);
102
103    /* Wait for all threads to complete */
104    for (i=0; i<NUM_THREADS; i++) {
105        pthread_join(threads[i], NULL);
106    }
107    printf ("Main(): Waited on %d  threads. Done.\n",
108 NUM_THREADS);
109
110    /* Clean up and exit */
111    pthread_attr_destroy(&attr);
112    pthread_mutex_destroy(&count_mutex);
113    pthread_cond_destroy(&count_threshold_cv);
114    pthread_exit(NULL);
115
116 }
```